KE5205 TEXT MINING 2017

TEXT CATEGORIZATION

Leong Mun Kew
Institute of Systems Science
National University of Singapore
email: munkew@nus.edu.sg

Objectives of this module

At the end of this module, you can:

- **Describe what is text categorization and how text categorization systems work**
- **Evaluate a text categorization system with respect to a business scenario**
- **Understand how supervised and unsupervised text categorization works**
- **Understand what is topic modeling**

Outline for this module

- **What is text categorization?**
- **How does supervised text categorization work?**
  - Document data set
  - Building a classifier
  - Evaluation
  - Running the classifier
- **Text categorization application examples**
- **Unsupervised text categorization**
  - Document clustering
- **Topic Modeling**

WHAT IS TEXT CATEGORIZATION?

Contrast with a library catalog

- **Example to right**
  - Subject: Statistics
- **Assigned by a cataloger**
- **Slow, tedious**
- **May be inconsistent**

Record 1 of 381 in National Library Board **Search type:** Search by Subjects
Search was: Statistics

Next

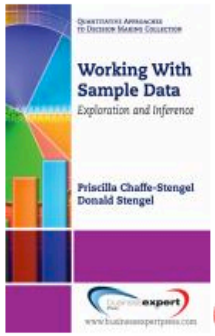

Annotation

Title Working with sample data : exploration and inference / Priscilla Chaffe-Stengel, Donald N. Stengel.
Author Chaffe-Stengel, Priscilla M.
Publisher New York : Business Expert Press, c2012.
Physical Description 151 p. : ill. ; 23 cm.
Notes Includes index.
"The quantitative approaches to decision making collection"--Cover.
Originally published in 2011.
Other Contributors Stengel, Donald N.
Search by Subjects Commercial statistics.
Statistics.

MESH index of a single journal paper

Below is an example of a **complete reference** in Medline (OvidSP) showing the journal article details and the list of MeSH headings (some with subheadings) assigned to it by the NLM Indexers:

Unique Identifier	20980007
Record Owner	From MEDLINE, a database of the U.S. National Library of Medicine.
Status	MEDLINE
Authors	Asejczyk-Widlicka M. Srodka W. Schachar RA. Pierscionek BK.
Authors Full Name	Asejczyk-Widlicka, M. Srodka, W. Schachar, R A. Pierscionek, B K.
Institution	Institute of Physics, Wroclaw University of Technology, Wybrzeze Wyspianskiego 27, 50-370 Wrocl Poland.
Title	Material properties of the cornea and sclera: a modelling approach to test experimental analysis
Source	Journal of Biomechanics. 44(3):543-6, 2011 Feb 3.
Abbreviated Source	J Biomech. 44(3):543-6, 2011 Feb 3.
NLM Journal Name	Journal of biomechanics
Publishing Model	Journal available in: Print-Electronic Citation processed from: Internet
NLM Journal Code	0157375, hjf
Country of Publication	United States
MeSH Subject Headings	Computer Simulation *Cornea / ph [Physiology] Finite Element Analysis Humans *Intraocular Pressure / ph [Physiology] Muscle Rigidity *Sclera / ph [Physiology] Visual Acuity / ph [Physiology]
Abstract	Material prope es of cornea and sclera are important for maintaining the shape of the eye and the requisite ce curvatures for optics. They also need to withstand the forces of external and inte sculature and fluctuations in intraocular pressure (IOP). These properties are difficult to iable results have been reported. A previously published experimental procedure, material properties of the eyeball coats were obtained, has been modelled in this te Element Analysis, in order to test the accuracy of the experiment. Material e calculated from the model and the resulting relationships between stress and rnea and sclera compared to their experimentally obtained counterparts. The ween model and experiment was close for the sclera but more varied for the cornea The pressure vessel model can be applied for measuring the material properties of the sclera but is less accurate for the cornea. Copyright Copyright 2010 Elsevier Ltd. All rights reserved.

Journal article details

MeSH descriptors assigned by indexers- taken from the list of preferred terms used to describe topics

The indexers have assigned the subheading Physiology to this MeSH descriptor

From: http://library.ulster.ac.uk/sci/mesh/

Automatic text categorization (also known as “classification”)

- **The process of assigning text documents uniquely into two or more categories (a document cannot be in more than one category)**
  - E.g., spam filtering – binary decision: “spam” or “not spam”

- **The process of assigning one or more category labels to a text document (a document may have more than one category)**
  - E.g., news filtering – which category to assign to news articles:
    - Sports, Olympics, Football (natural class)
    - Political, Business, Home,… (news sections)
    - Asian, Europe, Middle-East, …(geographical)

Some Examples of Text Classification

- Assigning subject categories to documents
- Email spam detection
- Medical diagnosis
- Identifying a language (before further processing)
- Identifying fraud (anomaly detection)
- Sentiment analysis (e.g., positive/negative reviews)
- Monitoring a news feed (e.g., news about Pope Francis)
- Etc.

ACCURACY

What is “accuracy”?

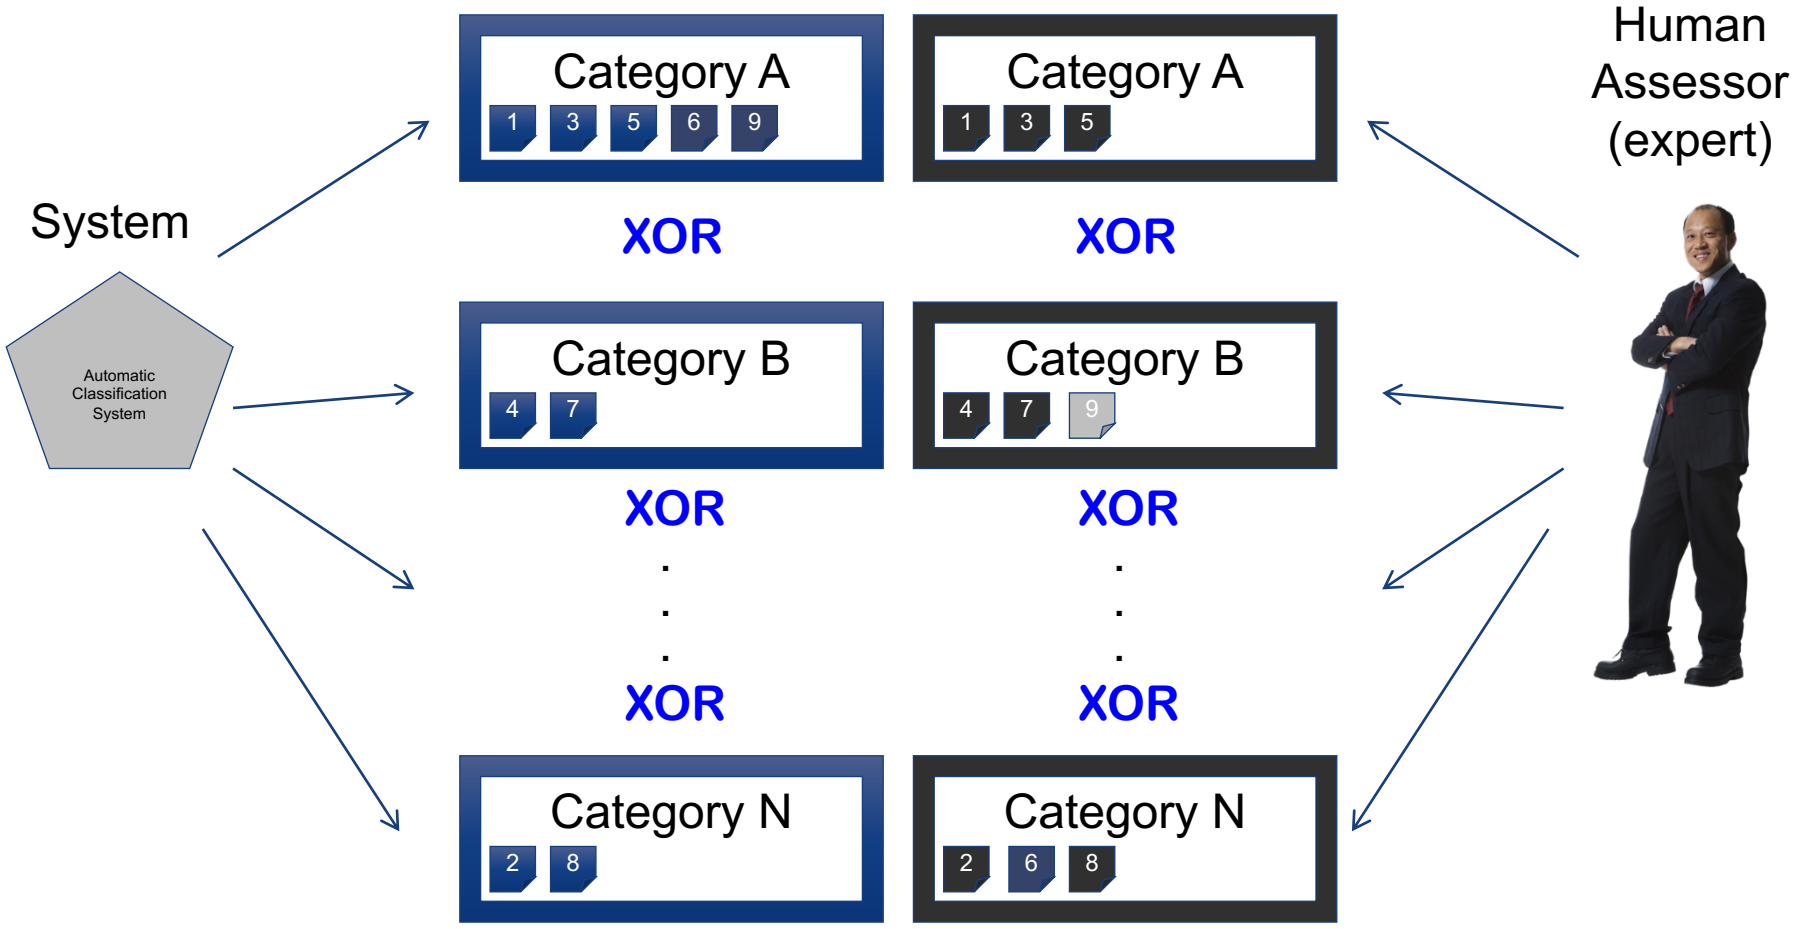

What do we mean by "Accuracy"

- **You measure an automatic categorization system by:**
  - How well it classifies a set of documents against a "reference"
  - This "reference" is normally a human expert
- **Reference**
  - Gold standard – accepted as being the best available
    - May not be perfect, e.g., tumour board for oncology
  - Good enough
    - Human expert(s), typically 80% agreement, good methodology
  - Better than nothing
    - "your boss tells you to do this, so you recruit your friends, family,…"
- **Most of the time, no such thing as "absolute truth"**

One measure of accuracy

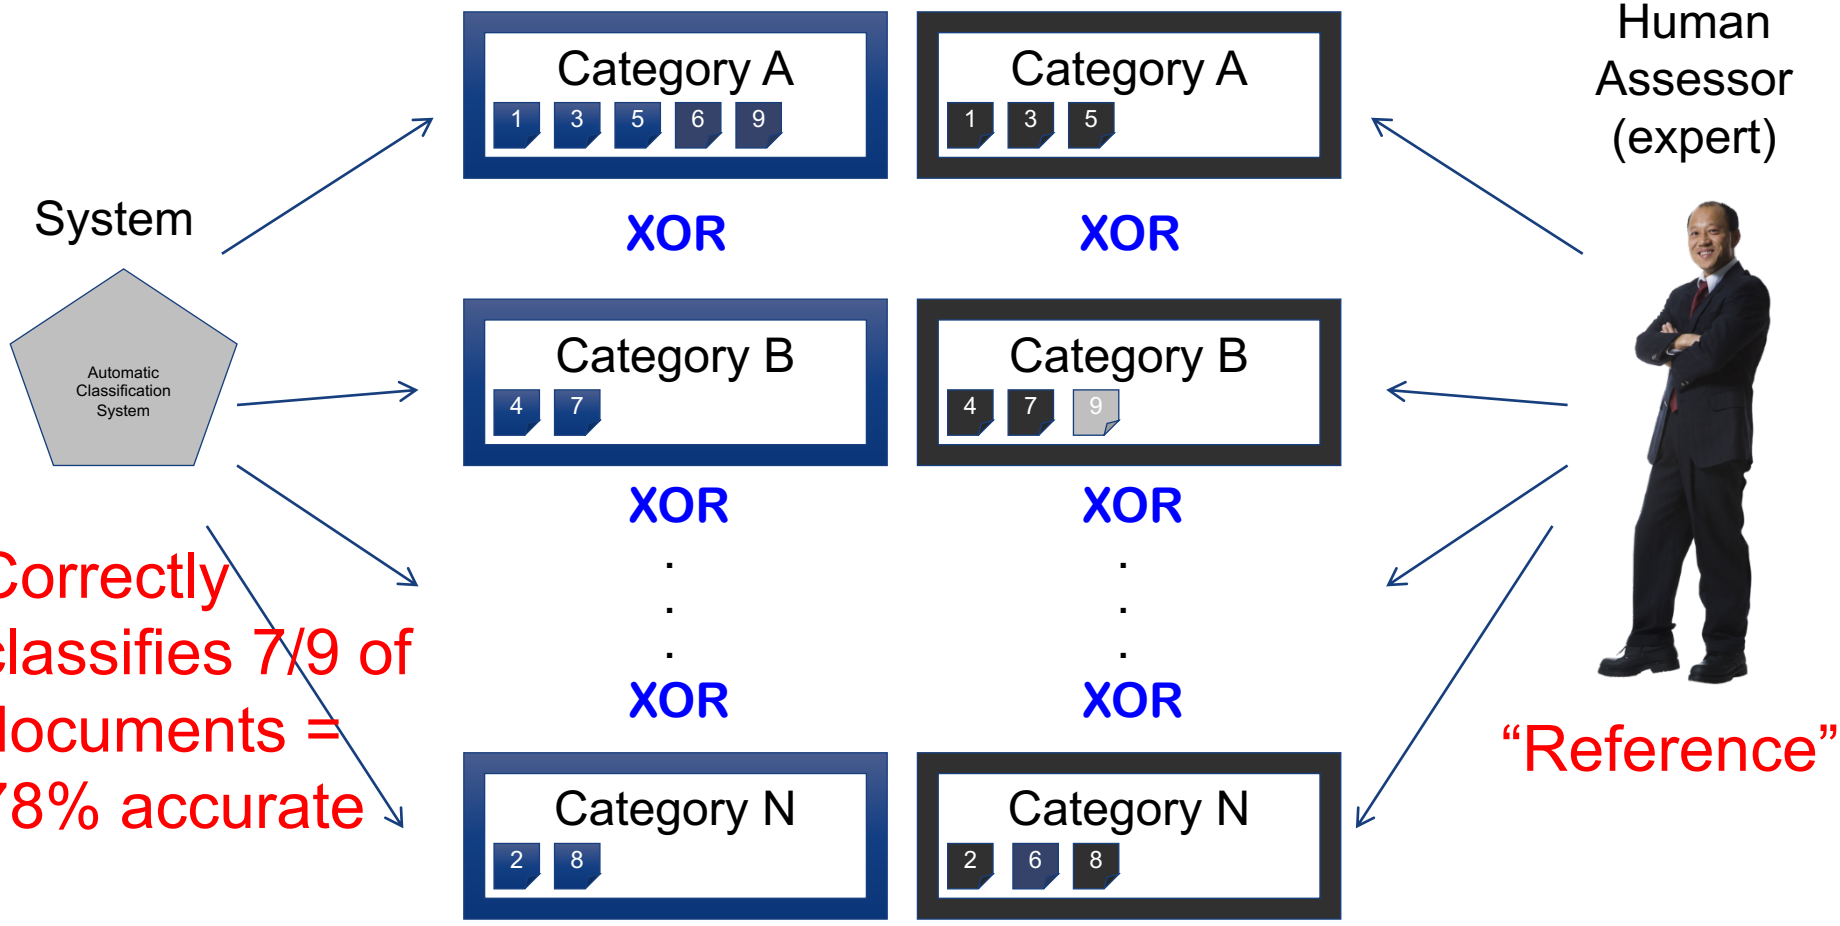

Discussion

- **Weather prediction (system predicts one week in advance)**

	Mon	Tue	Wed	Thu	Fri	Sat	Sun
System	Sunny	Drizzle	Rain	Sunny	Cloudy	Thunderstorms	Sunny
Actual	Sunny	Rain	Cloudy	Sunny	Drizzle	Thunderstorms	Cloudy

- **Questions:**
  - How “accurate” is the weather prediction system?
  - Do you have a tolerance for error? +/- margin of error?
  - How about outcomes – will I get wet if I use the system to decide whether or not to carry an umbrella? Will I get angry?

Discussion

- **There is a desert in USA where it only rains 10 days in a year. But when it rains, there are flash floods which kill an average of 20 people per year.**
- **You build a weather prediction system. How would you measure accuracy?**
  - If you said no rain every day, you would automatically be right 355/365 days in a year = 97% “accuracy”
  - If you said rain every day, you would be right only 10/365 = 3% of the time, but you would save 20 lives. Would this be true?
  - You are right half the time = 50%, so you would save 10 lives, but ONLY if people listen to your predictions.
- **What would you propose?**

HOW DOES AUTOMATIC TEXT CATEGORIZATION WORK?

Text Categorization Phases

Two Phases

1. **Training – creating the text "classifier" (automatic categorization engine)**
   - You need a set of documents, already categorized
   - Divide the set into training (typically 70%) and testing (30%)
   - Build your classifier such that it's able to accurately classify the training set of documents to your level of comfort
     - "level of comfort" depends on how hard is the task! ☺
   - Evaluate your classifier on the test set; ensure sufficient accuracy
2. **Running – using your classifier on new sets of documents**
   - You will not know how well it performs
   - Need to "audit" the results occasionally (use an assessor)
     - Assess random sample of the documents against the predicted categories

DOCUMENT DATA SET

Movie reviews classified as “good” and “bad”

POSITIVE POLARITY (GOOD)

-a mesmerizing cinematic poem from the first frame to the last .”
-a well-put-together piece of urban satire .
-one can't deny its seriousness and quality .
-hard to resist .
-a naturally funny film , home movie makes you crave chris smith's next movie .
-a true-blue delight .
-a fun ride .
-a surprisingly funny movie .
-the script is smart and dark - hallelujah for small favors .
-a flick about our infantilized culture that isn't entirely infantile .

-unfortunately the story and the actors are served with a hack script .
-too slow for a younger crowd , too shallow for an older one .
-terminally brain dead production .
-one lousy movie .
-this movie . . . doesn't deserve the energy it takes to describe how bad it is .
-a cleverly crafted but ultimately hollow mockumentary .
-it's an 88-minute highlight reel that's 86 minutes too long .
-the whole affair is as predictable as can be .

NEGATIVE POLARITY (BAD)

From: http://karpathy.ca/mlsite/lecture2.php

Movie reviews classified as “good” and “bad”

POSITIVE POLARITY (GOOD)

-a mesmerizing cinematic poem from the first frame to the last .”
-a well-put-together piece of urban satire .
-one can't deny its seriousness and quality .
-hard to resist .
-a naturally funny film , home movie makes you crave chris smith's next movie .
-a true-blue delight .
-a fun ride .
-a surprisingly funny movie .
-the script is smart and dark - hallelujah for small favors .
-a flick about our infantilized culture that isn't entirely infantile .

5000 reviews

-unfortunately the story and the actors are served with a hack script .
-too slow for a younger crowd , too shallow for an older one .
-terminally brain dead production .
-one lousy movie .
-this movie ... doesn't deserve the energy it takes to describe how bad it is .
-a cleverly crafted but ultimately hollow mockumentary .
-it's an 88-minute highlight reel that's 86 minutes too long .
-the whole affair is as predictable as can be .

NEGATIVE POLARITY (BAD)

5000 reviews

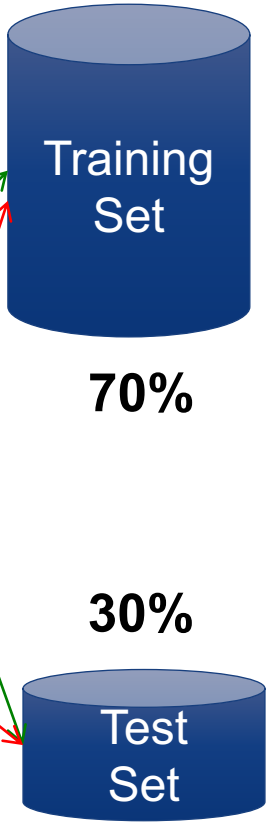

From: http://karpathy.ca/mlsite/lecture2.php

BUILDING A CLASSIFIER

JUST SOME EXAMPLES (NOT EXHAUSTIVE)

Creating classifiers

- **Hand-coded classifiers (the "good old days!")**
  - If <conditions> then <category> else NOT<category>, where conditions are normally in disjunctive normal form

If	((wheat & farm)	**or**	
	(wheat & commodity)	**or**	
	(bushels & export)	**or**	
	(wheat & tonnes)	**or**	
	(wheat & winter & ¬soft))	**then**	WHEAT else ¬ WHEAT

From: F. Aiolli, *Text Categorization*, http://www.math.unipd.it/~aiolli/corsi/SI-0607/Lez09.251006.pd

Inducing classifiers (1)

- **Probabilistic Classifiers**
  - Represent the probability that a document d_i belongs to category c_j by

    $$P(c_j|d_i) = P(c_j)P(d_i|c_j)/P(d_i)$$

    where $P(c_j)$ is the probability that a randomly selected document belongs to category c_j, $P(d_i)$ is the probability that a randomly selected document has d_i as a vector representation and $P(d_i|c_j)$ is the probability that category c_j contains d_i.
- **Decision Tree Classifiers**
  - Uses symbolic representation in a tree network
  - Internal nodes are terms, edges are tests on the weights, and leaf nodes are categories
  - The classifier works by traversing a path to the appropriate leaf node

Exercise

- **List of boys names**
  - Alan
  - Barry
  - Colin
  - Dexter
  - Edward
  - Frederick
  - Howard
  - ….
- **List of girls names**
  - Anna
  - Betty
  - Chelsea
  - Doris
  - Elizabeth
  - Fanny
  - Hortense
  - ….

Given a list of boys' names and a list of girls' names, build a simple classifier to distinguish boys' names from girls' names

Example of a decision tree to decide if a name is male or female

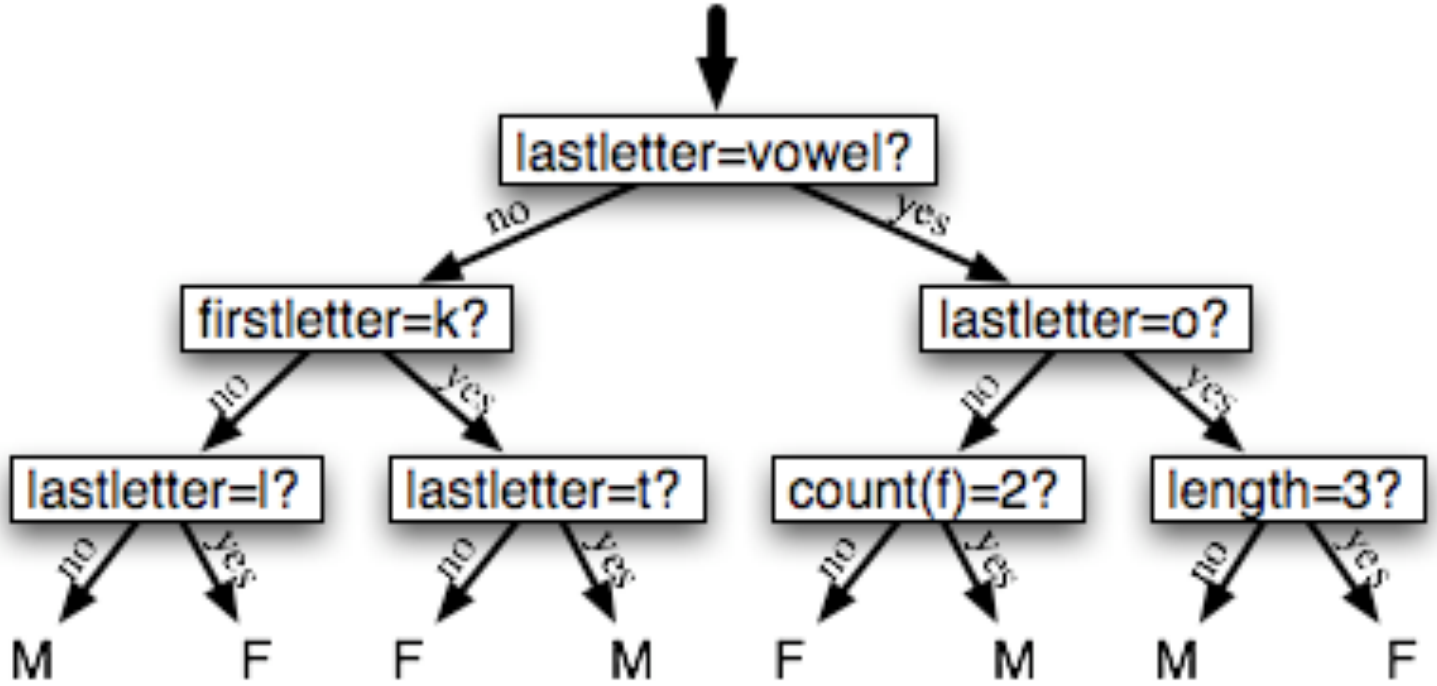

From: http://nltk.googlecode.com/svn/trunk/doc/book/ch06.html

Inducing classifiers (2)

- **The Rocchio Classifiers**
  - Each category is represented by a prototypical document, i.e., profile vector
  - Documents are classified by similarity to the profile vector

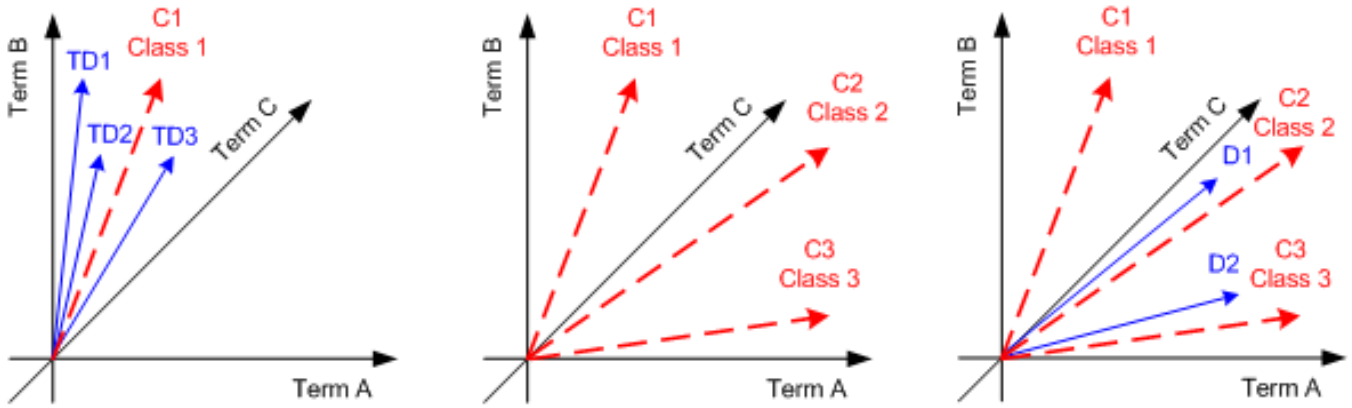

A) Training Document Representations in Vector Space define Class Representation (Centroid)

B) Class Representation by Class Vectors (CentroidS)

C) Classification of Documents by similarity between Document Vector and Class Vector

From: http://www.iicm.tugraz.at/about/Homepages/cguetl/courses/isr/opt/classification/Vector_Space_Model.html

Inducing classifiers (3)

- **Support Vector Machines (SVMs)**
  - SVMs divide the term space in hyperplanes separating the positive and negative training samples.
  - The surface that provides the widest separation between the support surfaces is selected

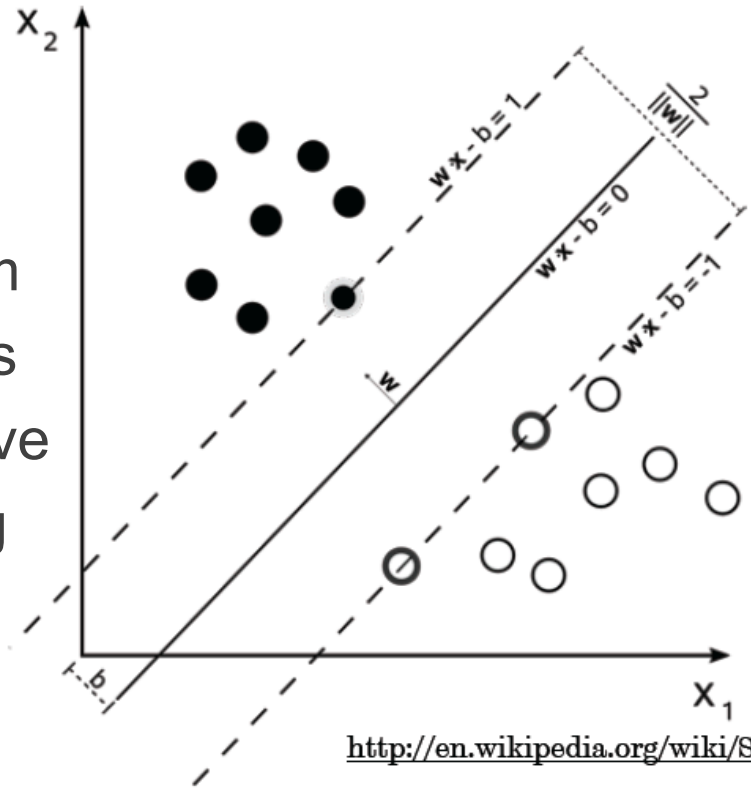

http://en.wikipedia.org/wiki/Support_vector_machines

EVALUATION

COMPARING DIFFERENT CLASSIFIERS FOR THE SAME CATEGORIZATION TASK

Remember this?

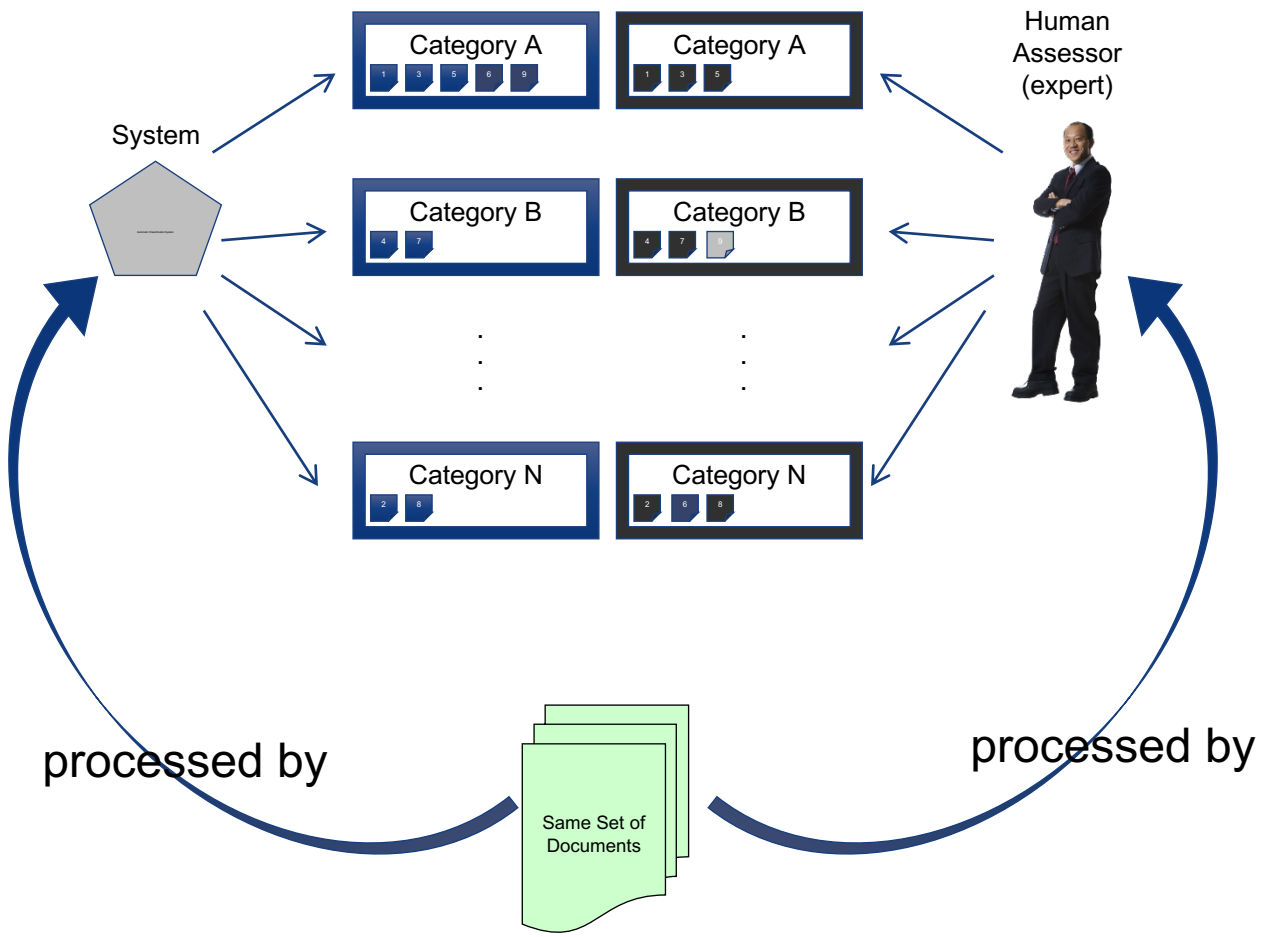

Confusion Matrix

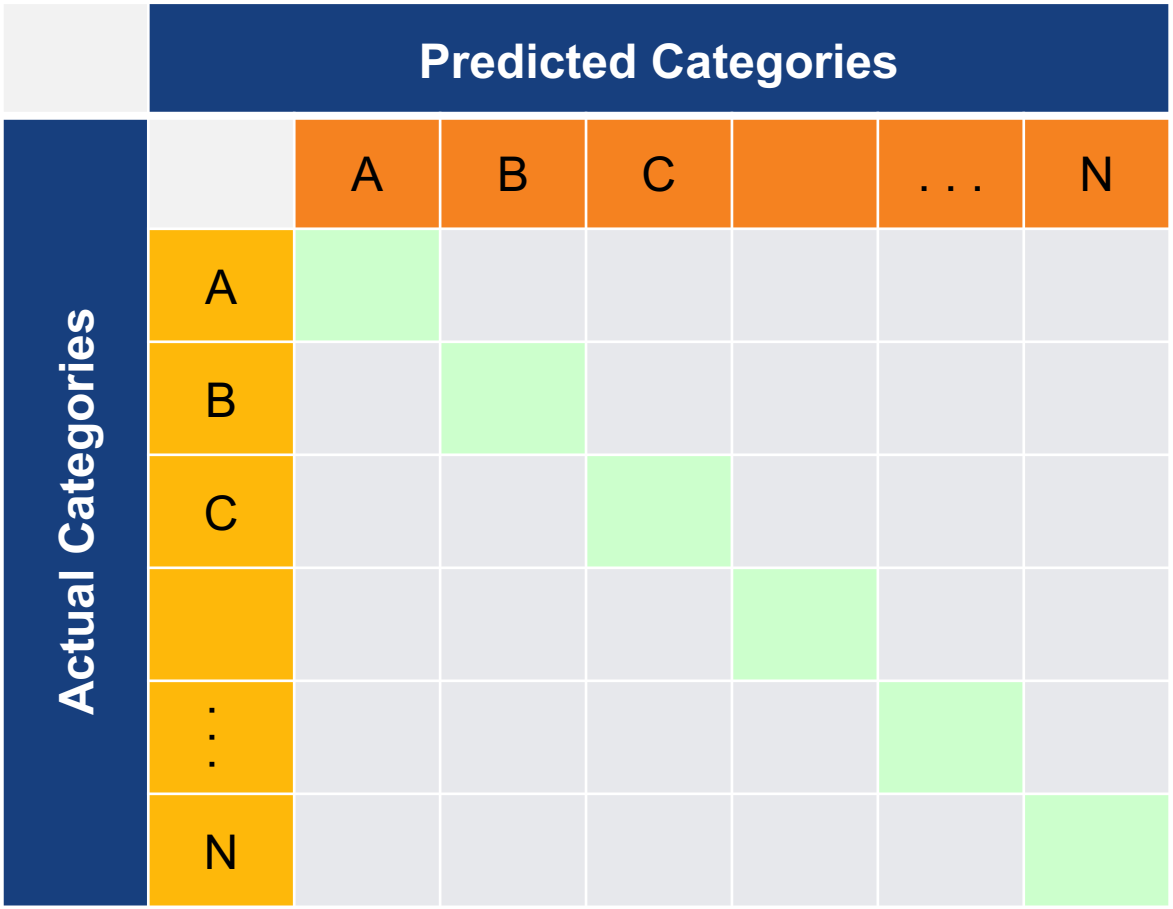

Example (using %)

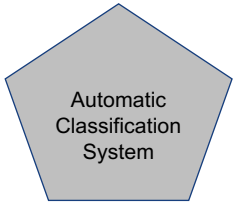

Actual Categories \ Predicted Categories	A	B	C		. . .	N	
A	87%	2%	5%		. . .	1%	= 100%
B	6%	90%	0%		. . .	2%	= 100%
C	12%	2%	77%		. . .	4%	= 100%
					. . .		
⋮					. . .		
N	21%	0%	4%		. . .	65%	= 100%

Example (using #)

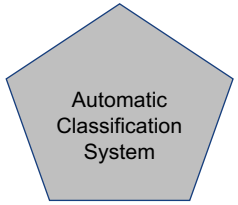

		Predicted Categories						
			A	B	C		. . .	N	
Actual Categories	A		143	34	17		. . .	2	= Tot(A) docs
	B		67	1289	44		. . .	239	= Tot(B) docs
	C		980	234	3454		. . .	88	= Tot(C) docs
							. . .		
	⋮						. . .		
	N		87	24	63		. . .	650	= Tot(N) docs

Consider the simple 2x2 matrix (2000 documents were classified)

Desired positive prediction

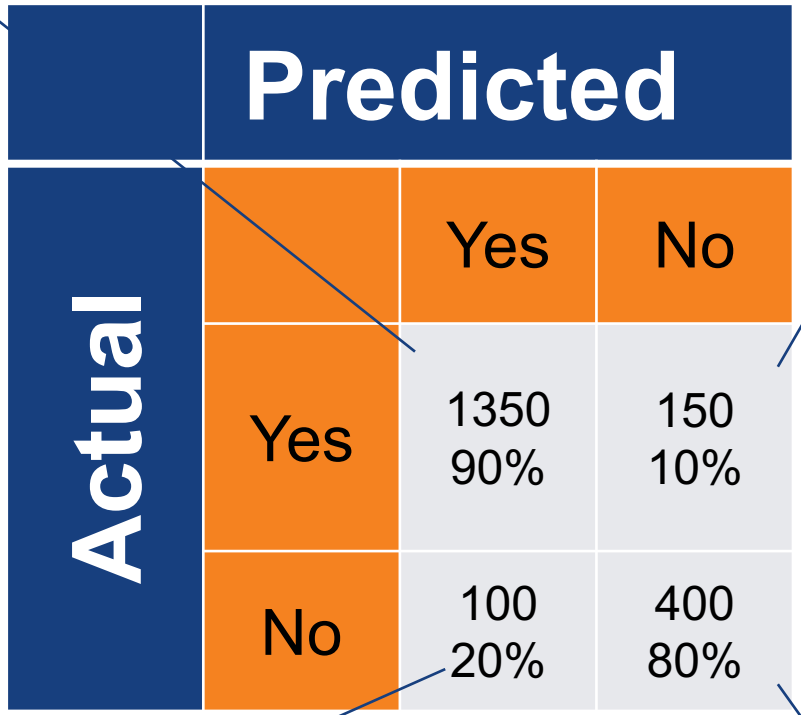

False negative

False positive

Desired negative prediction

EVALUATING MULTIPLE CLASSIFIERS

Discussion

- **Weather prediction. You ask two people to predict whether it will rain or not in the coming week:**

	Mon	Tue	Wed	Thu	Fri	Sat	Sun
ZZ	Yes	No	Yes	No	Yes	No	No
MK	No	No	No	Yes	No	No	Yes
Actual	Yes	Yes	Yes	No	Yes	Yes	No

- **Questions:**
  - Who is more “accurate”? ZZ is right 5/7 times. MK is right 0/7 times.
  - Who should you ask in future if you don’t want to get wet?
  - Who is the better “classifier”?

What happens with 2 classifiers?

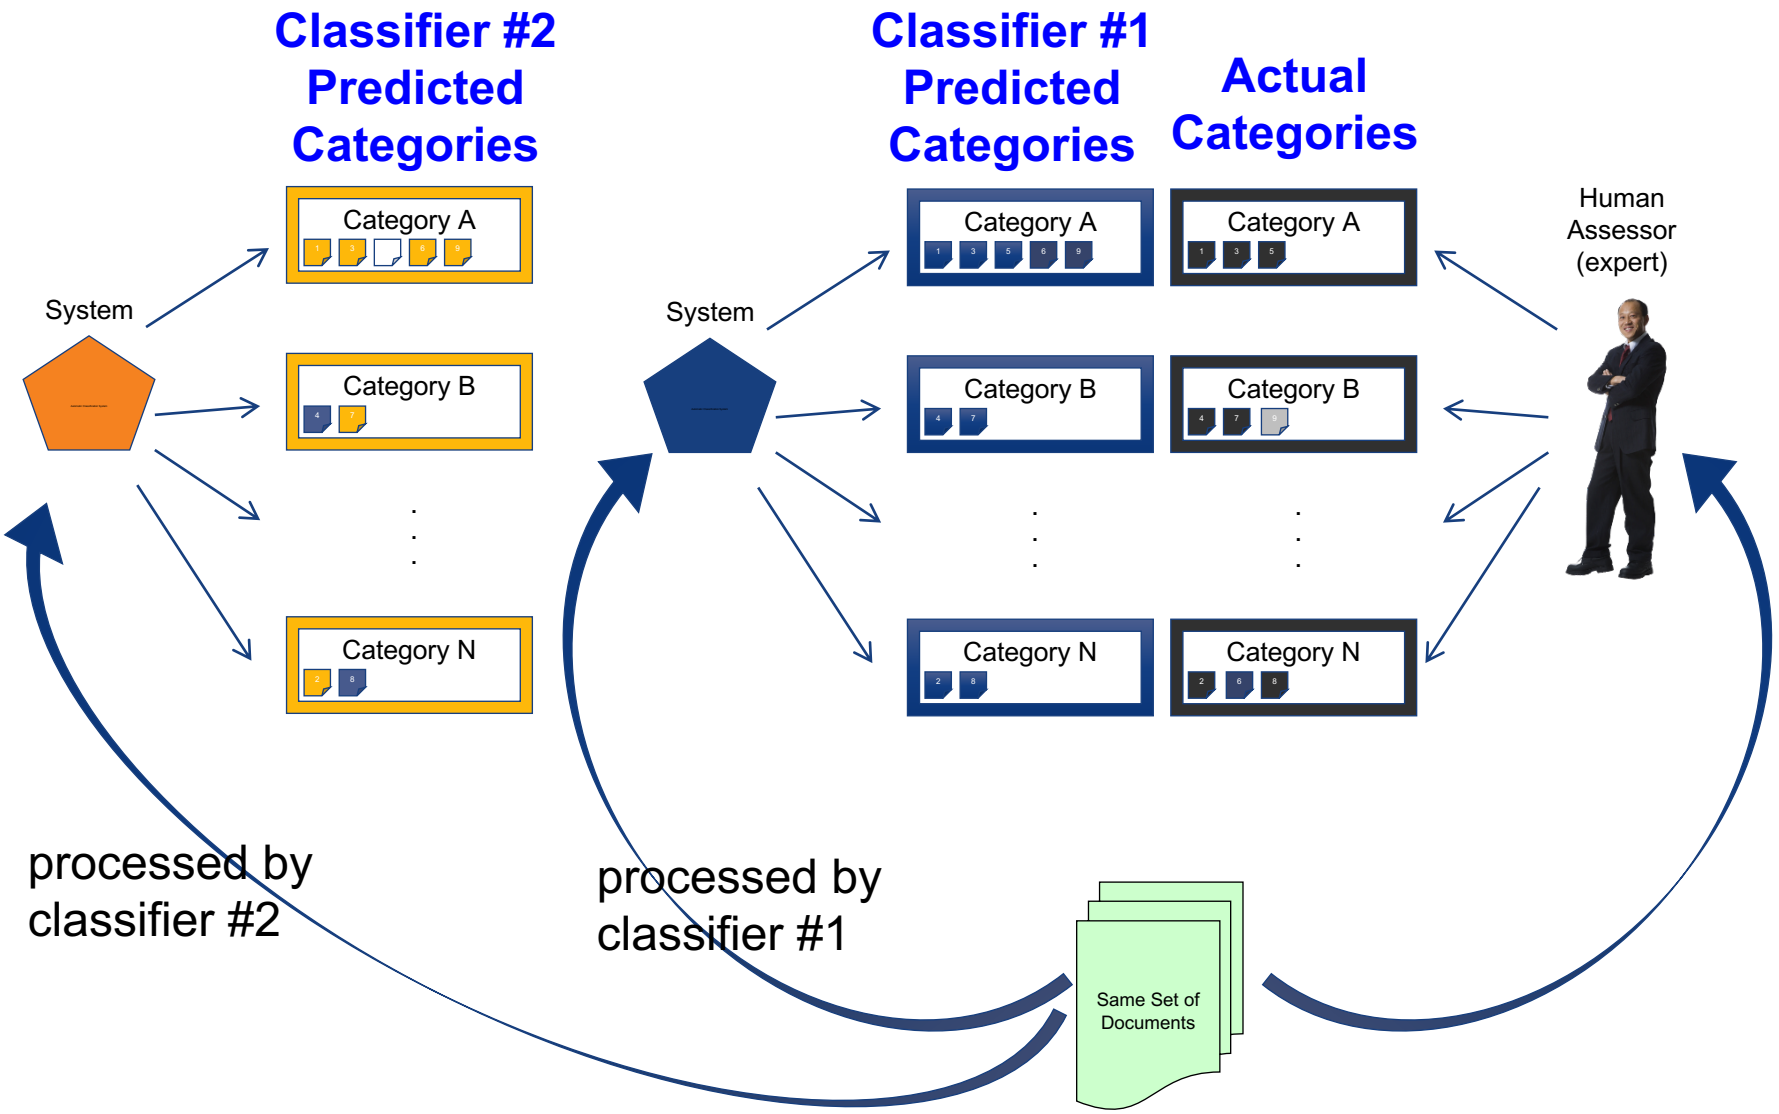

Comparing models: e.g. 2x2 matrix (actual numbers of documents)

Classifier #1

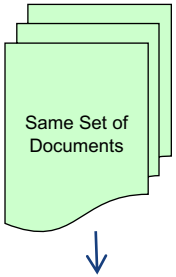

Classifier #2

		Predicted	
		Y	N
Actual	Y	900	100
	N	40	410

Seems quite good for both predictions

		Predicted	
		Y	N
Actual	Y	700	300
	N	2	448

Reduced the false positives but false negatives increased

? Which classifier is better?

Adding the semantics – the courtroom

Classifier #1

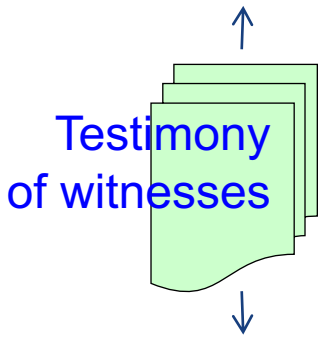

Classifier #2

Actual	Predicted		
		Guilty	Innocent
	Guilty	900	100
	Innocent	40	410

Let 100 guilty go free
Convict 40 innocent persons

? Which classifier is better?

Actual	Predicted		
		Guilty	Innocent
	Guilty	700	300
	Innocent	2	448

Let 300 guilty go free
Convict 2 innocent persons

Adding a cost function – fraud investigation

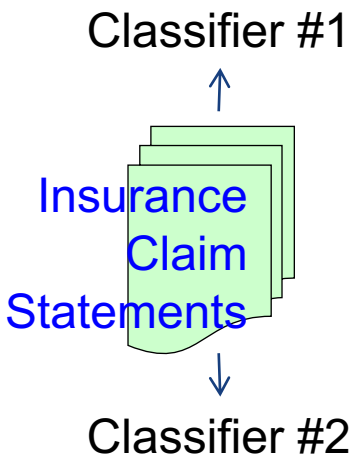

Actual \ Predicted	Genuine	Fraud
Genuine	900	100
Fraud	40	410

Company loses $80k in fraud
Company pays $255k in costs

? Which classifier is better?

Actual \ Predicted	Genuine	Fraud
Genuine	700	300
Fraud	2	448

Company loses $4k in fraud
Company pays $374k in costs

The average fraud costs the company $2000
It costs the company $500 to investigate each suspected fraud

Adding a cost function – fraud investigation

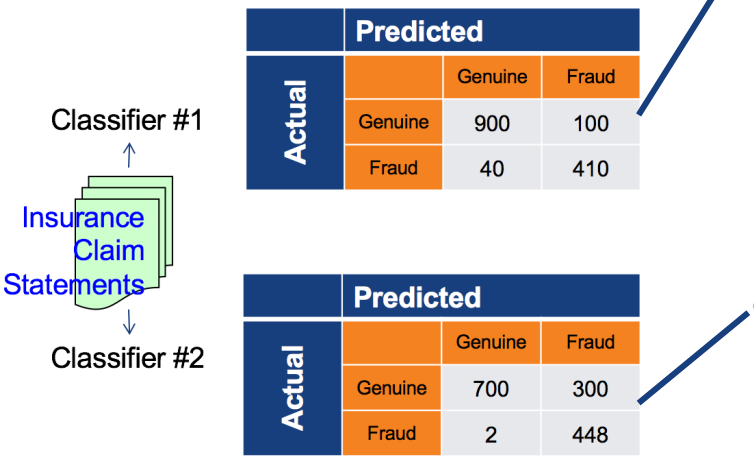

- ***Consider Doing nothing (don't act to identify fraud):***
  - *Predicted fraud = 0 cases @$500 per case costs $0k for investigation.*
  - *Undetected fraud is 450 cases @$2k/fraud loses $900k.*
  - *Overall -$0k -$900k = -$900k*
- ***Analysis for classifier #1:***
  - *Predicted fraud = 510 cases @$500 per case costs $255k for investigation.*
  - *Undetected fraud is 40 cases @$2k/fraud loses $80k.*
  - *Overall -$255k -$80k = -$335k*
- ***Analysis for classifier #2:***
  - *Predicted fraud = 748 cases @$500 per case costs $374k for investigation.*
  - *Undetected fraud is 2 cases @$2k/fraud loses $4k.*
  - *Overall -$374k -$4k = -$378k*

Adding a cost function – fraud investigation

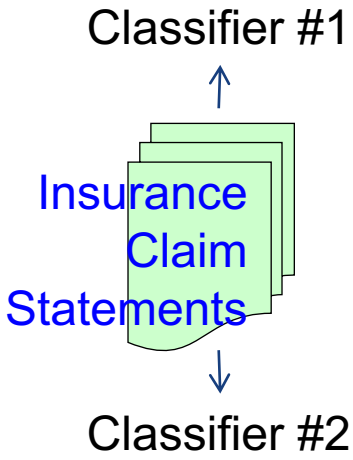

Actual	Predicted: Genuine	Predicted: Fraud
Genuine	900	100
Fraud	40	410

Company loses $80k in fraud
Company pays $255k in costs

Company paid $119k to save $76k

Actual	Predicted: Genuine	Predicted: Fraud
Genuine	700	300
Fraud	2	448

Company loses $4k in fraud
Company pays $374k in costs

The average fraud costs the company $2000
It costs the company $500 to investigate each suspected fraud

Consider the branding?

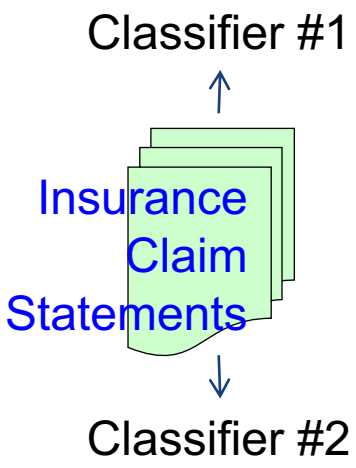

Actual	Predicted	
	Genuine	Fraud
Genuine	900	100
Fraud	40	410

Company loses $80k in fraud
Company pays $255k in costs

Company paid $119k to save $76k

Actual	Predicted	
	Genuine	Fraud
Genuine	700	300
Fraud	2	448

Company loses $4k in fraud
Company pays $374k in costs

The average fraud costs the company $2000
It costs the company $500 to investigate each suspected fraud

Classifier evaluation

- **Evaluation of classifiers is done with respect to a *business context***
- **Evaluation of classifiers is normally done *empirically***
- **Experimental evaluation focuses on *effectiveness*, i.e., the ability of the classifier to make the right classification decision**
- **Precision & Recall concepts as applied to (multi-class) categorization**
  - Precision is the probability that if a random document d_i is categorized under category c_j, that decision is correct
  - Recall wrt c_j is the probability that if a random document d_i should be categorized under c_j , then the decision is taken

RUNNING THE CLASSIFIER

Expect False Results
eg: spam filtering

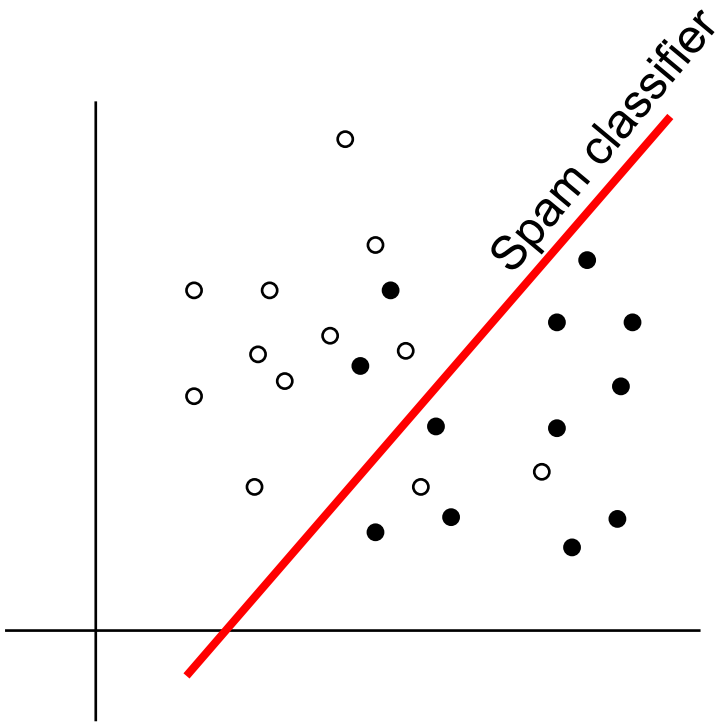

- Email data – non-spam
- Email data – spam

Training Set

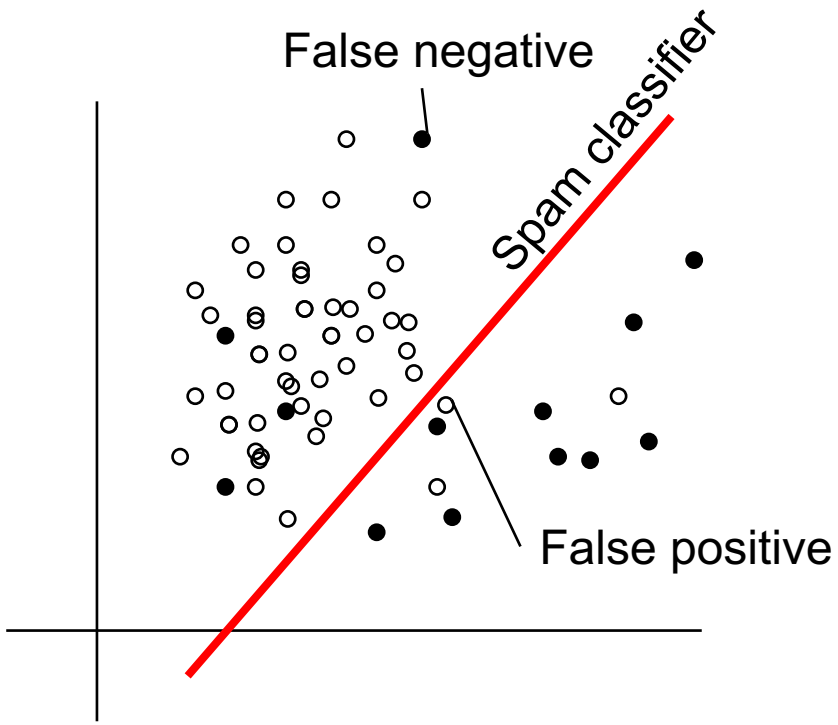

- Email non-spam
- Email spam

Real email stream

Overfitting the Training Set

Training Set

Actual \ Predicted	Y	N
Y	700	0
N	0	300

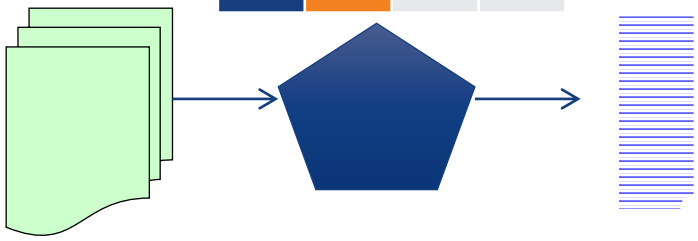

Training Set

Actual \ Predicted	Y	N
Y	600	100
N	50	250

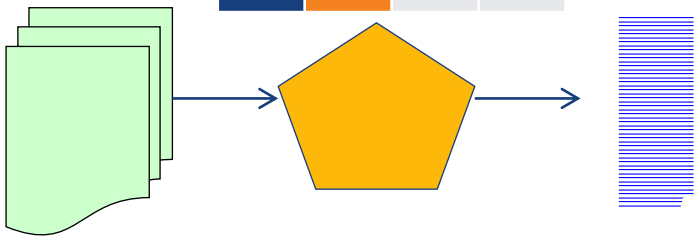

? Which classifier is better?

Overfitting the Training Set

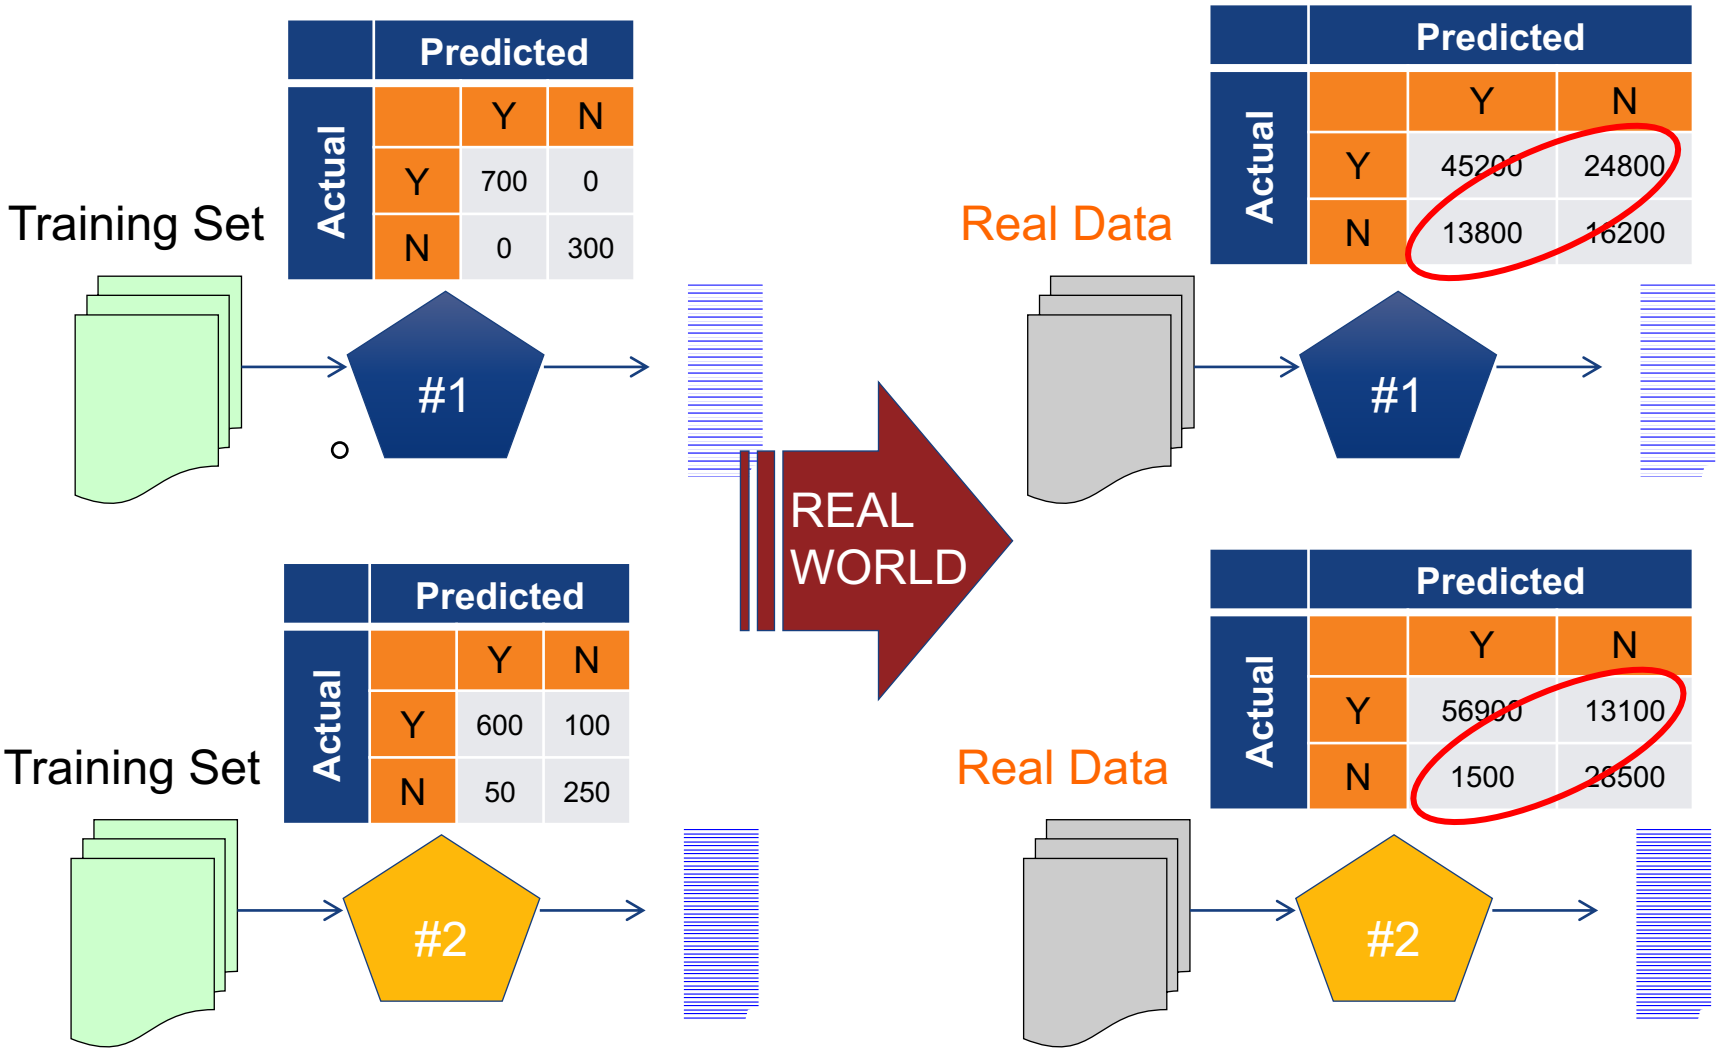

Overfitting the Training Set – what happened?

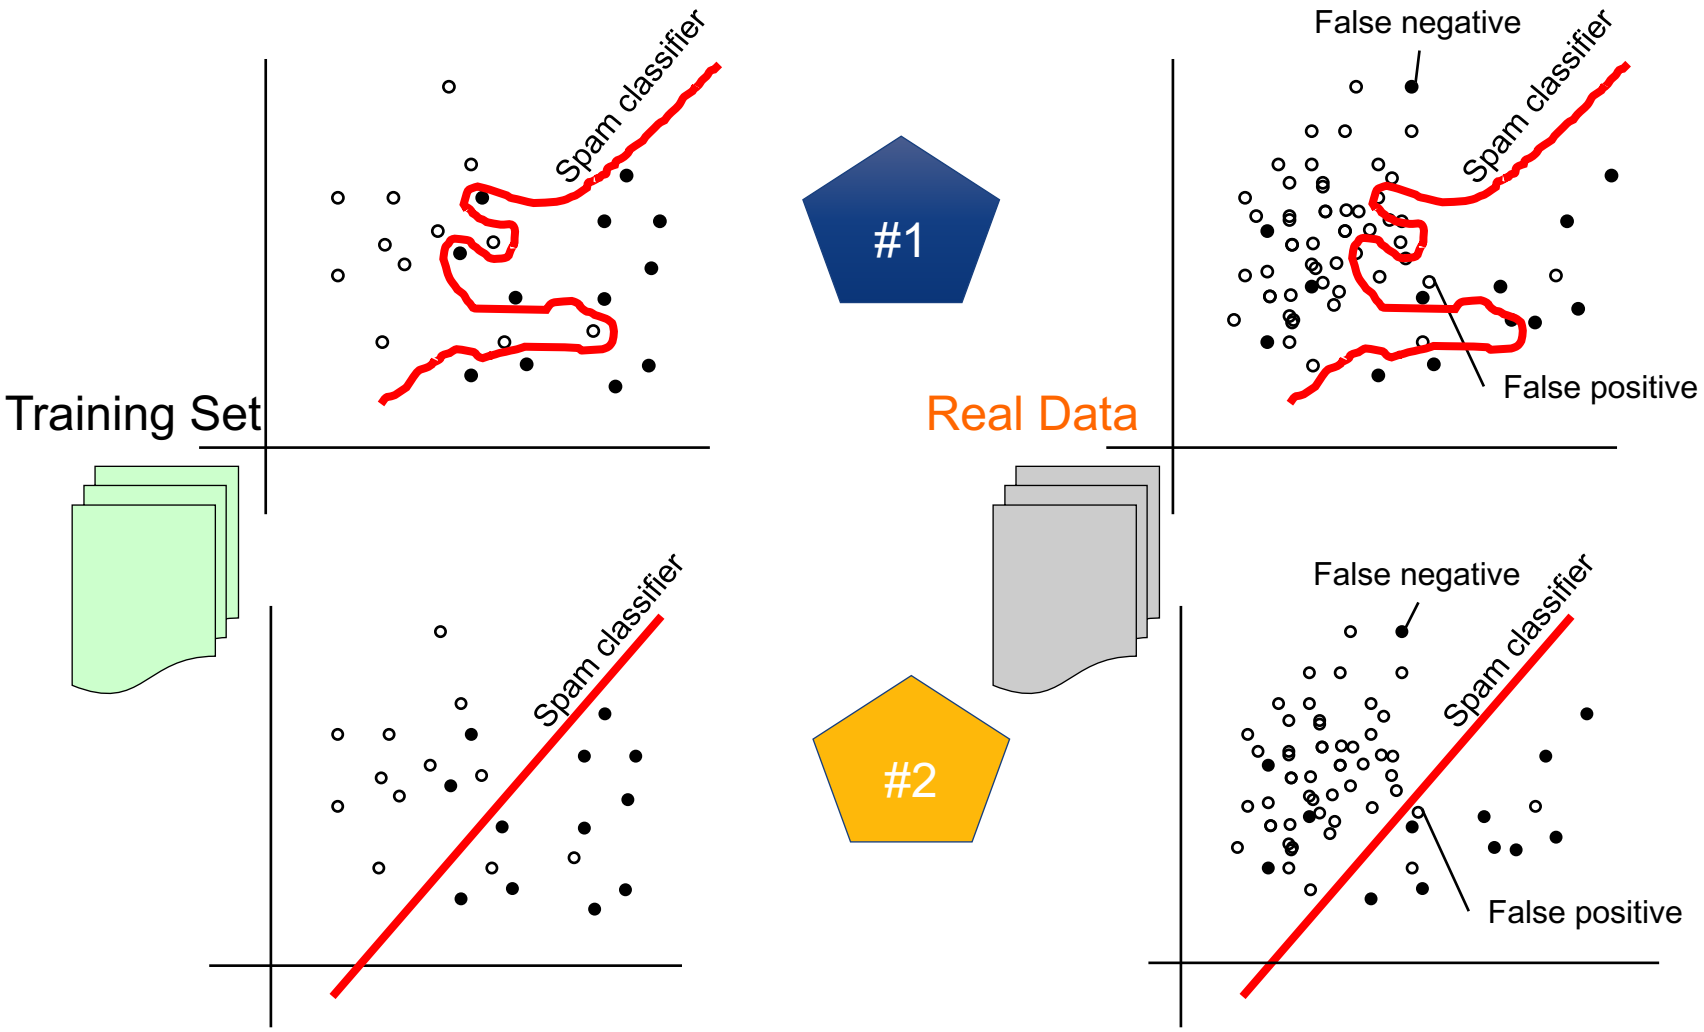

Hard and Soft Categorization

- **Fully automated classifiers make "hard" binary decisions**
  - In example to right, the document, D, is assigned to category C only.

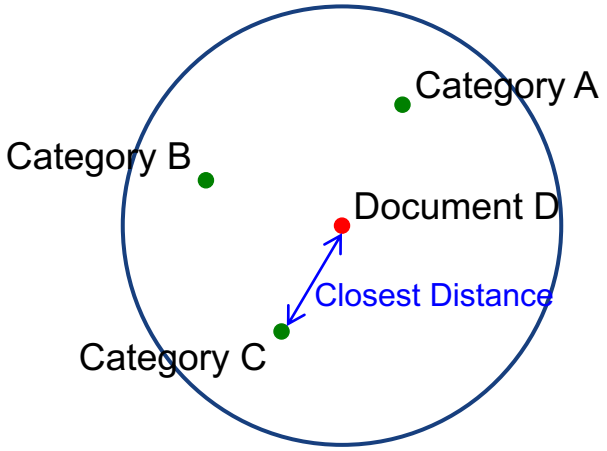

- **Semi-automated (interactive) classifiers instead are created by allowing "soft" real-value decisions**
  - Rank the categories according to their measure of appropriateness for the document
  - In example to right, the document, E, is assigned 3 possible categories:

Rank	Category	Probability
1	Z	0.76
2	X	0.72
3	Y	0.54

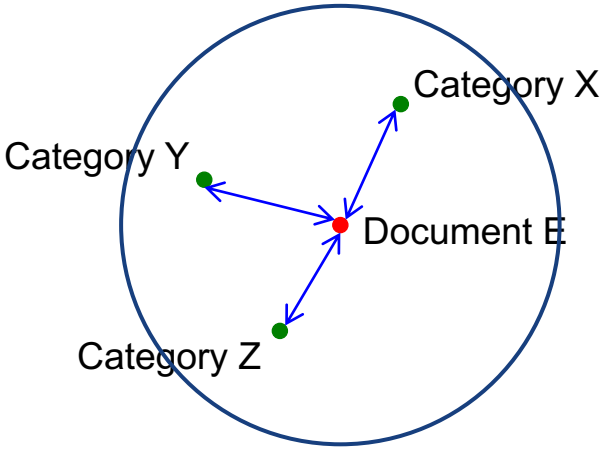

  - Used for computer assisted human decision making
    - For example, in critical applications such as medical diagnosis

Aggregating multiple classifiers

Running with more than one classifier

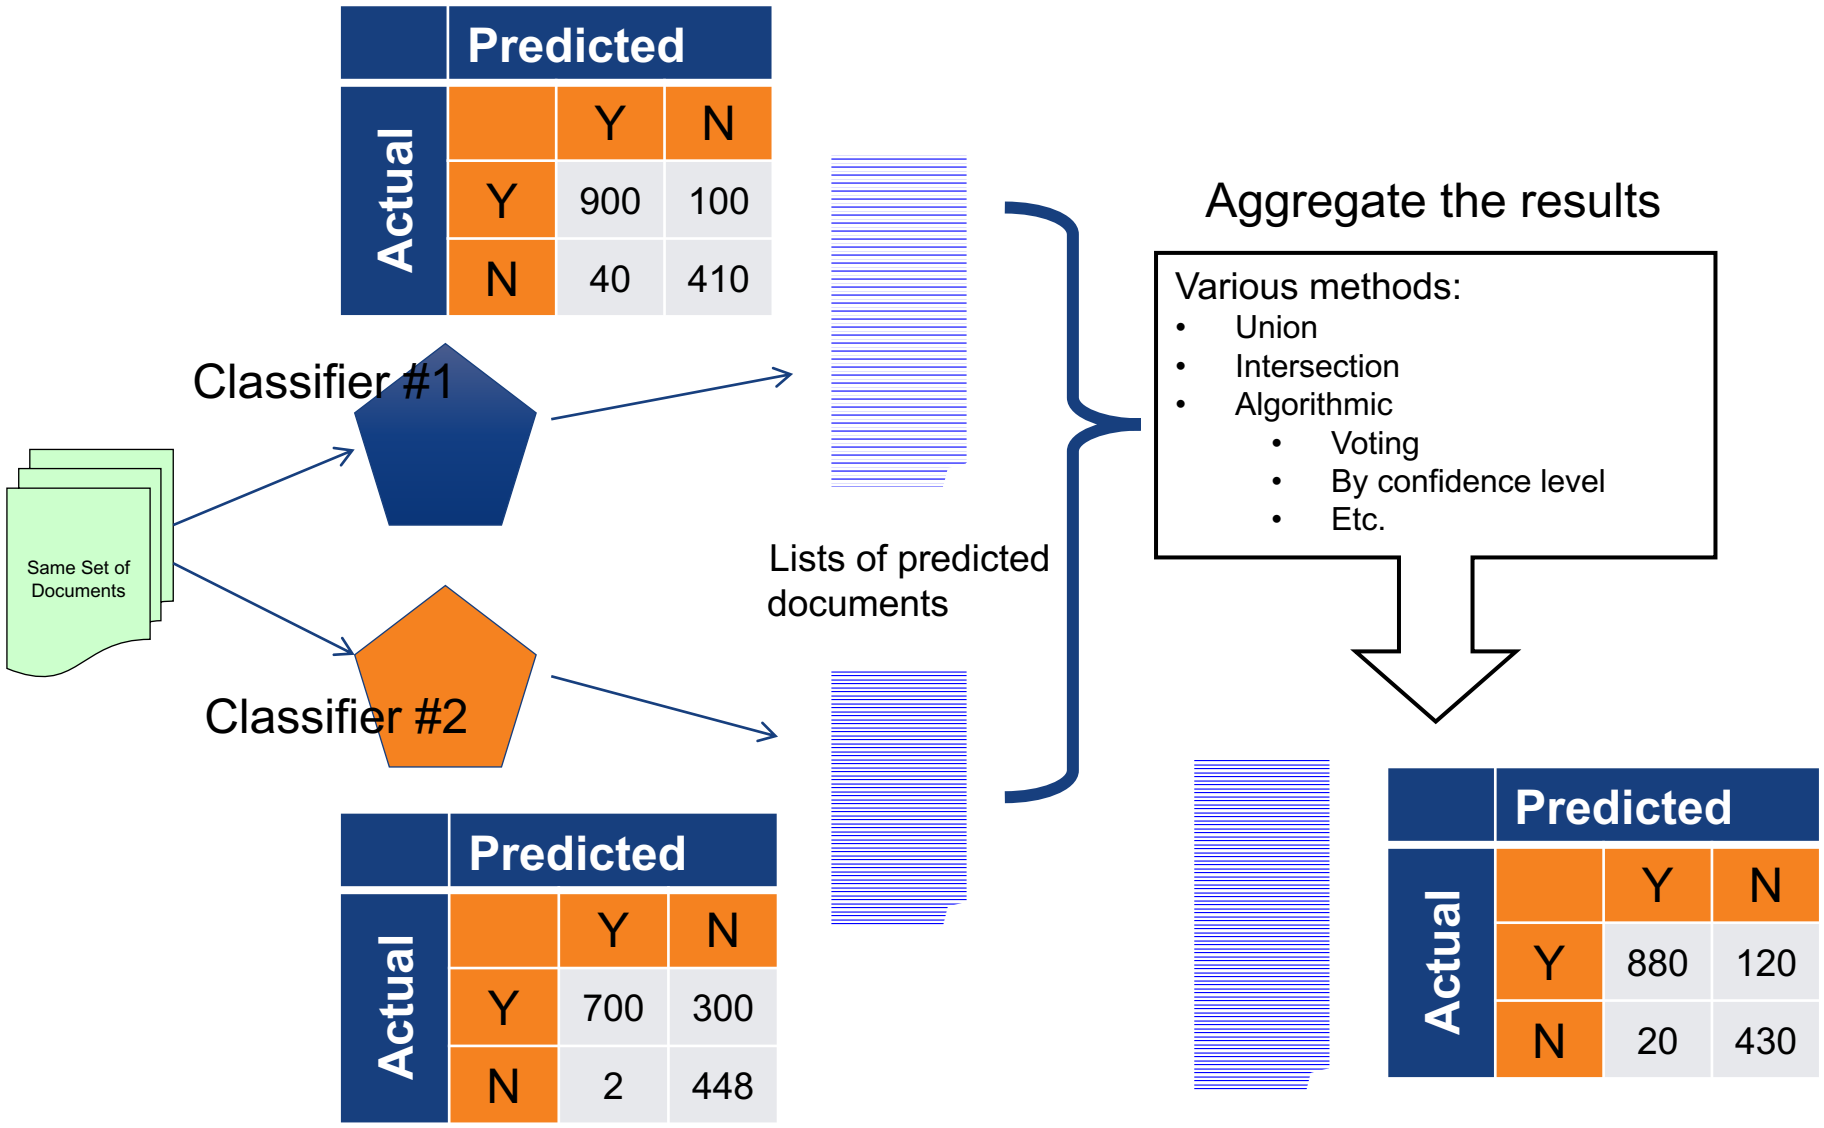

TEXT CATEGORIZATION APPLICATION EXAMPLES

Boosting Identification of Fraudulent Claims

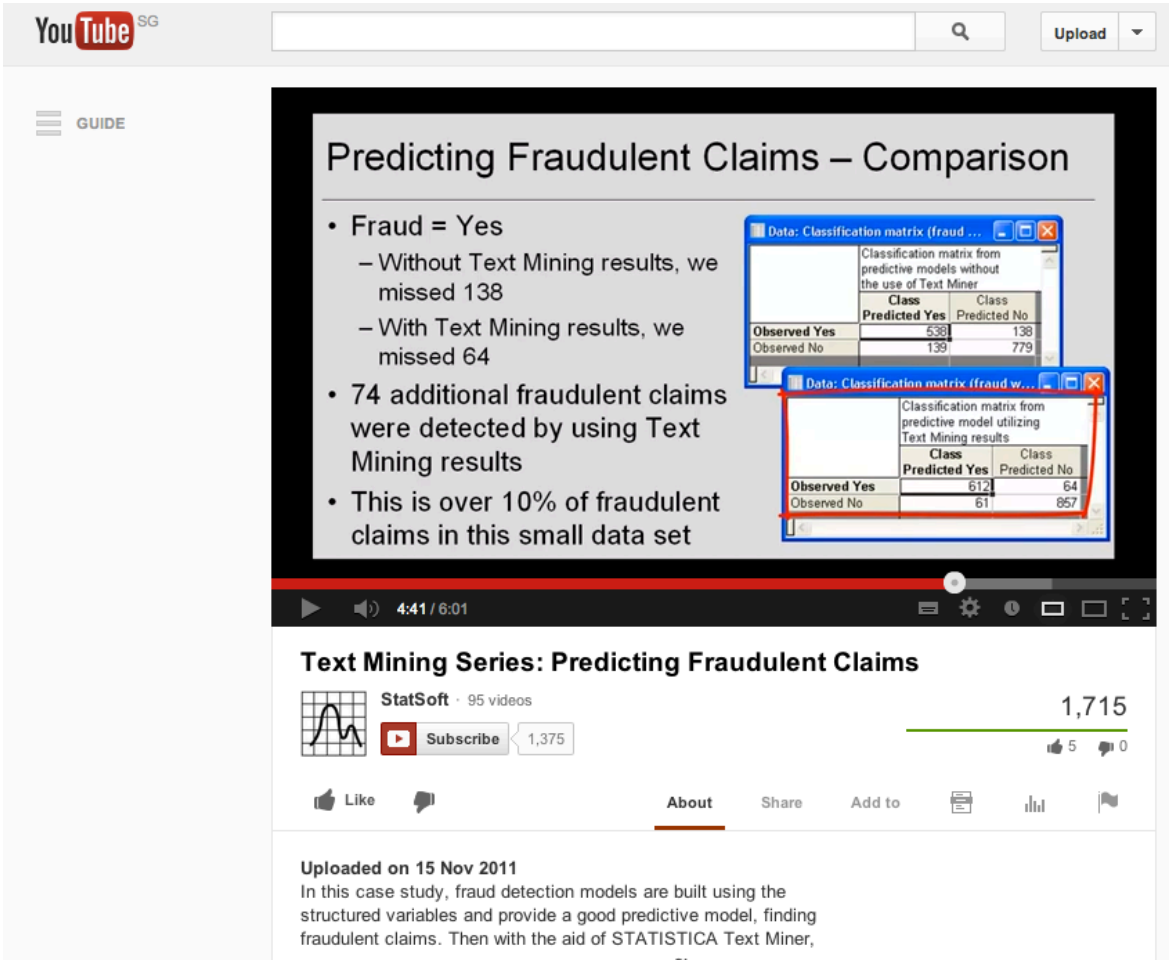

From: http://www.youtube.com/watch?v=OlQpm8qTog4

Automatic Categorization of Documents

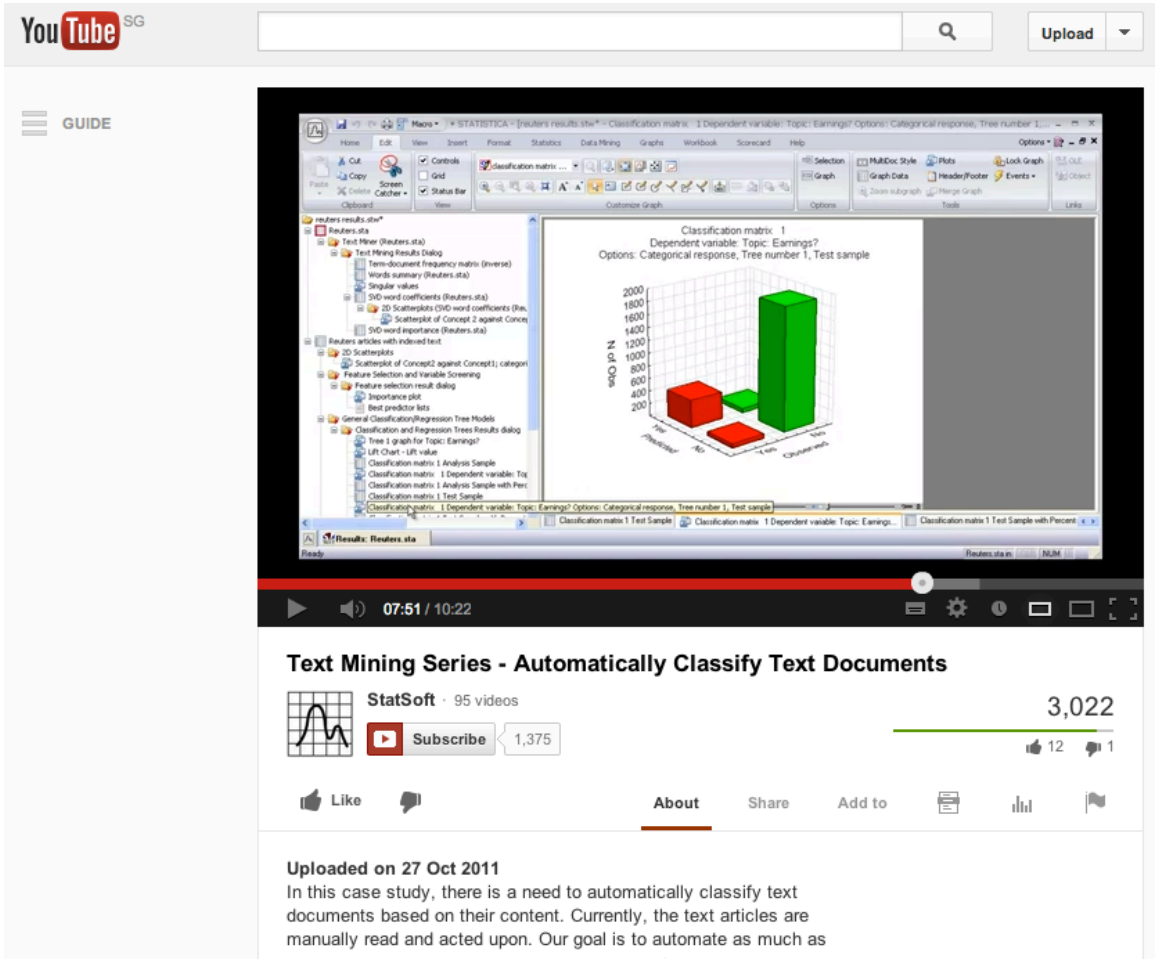

From: http://www.youtube.com/watch?v=Q5K3gyQJkC0

UNSUPERVISED TEXT CATEGORIZATION

Supervised vs Unsupervised

- **Supervised categorization**

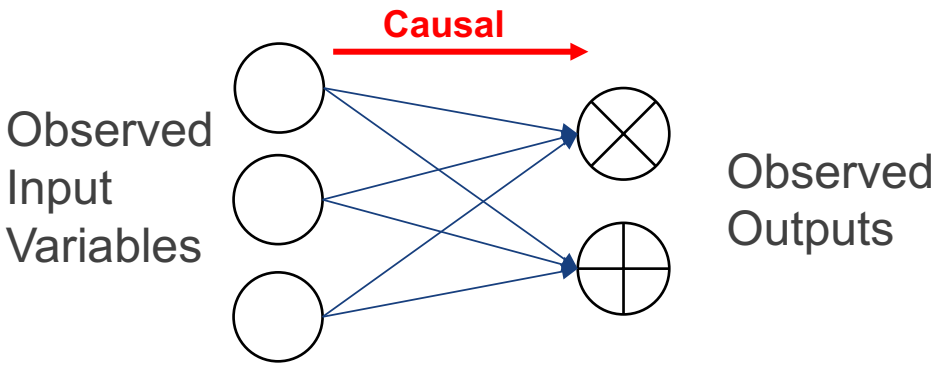

Essentially similar models
- **Inputs causally effects the outputs**
- **Input variables may not be observable**

- **Unsupervised categorization**

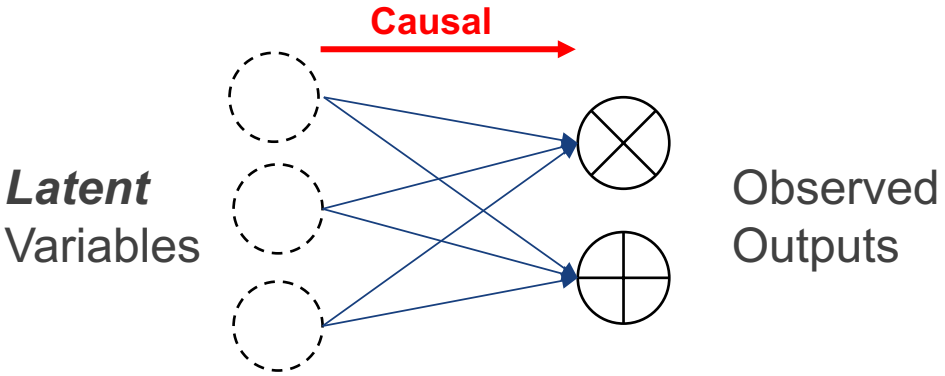

Example

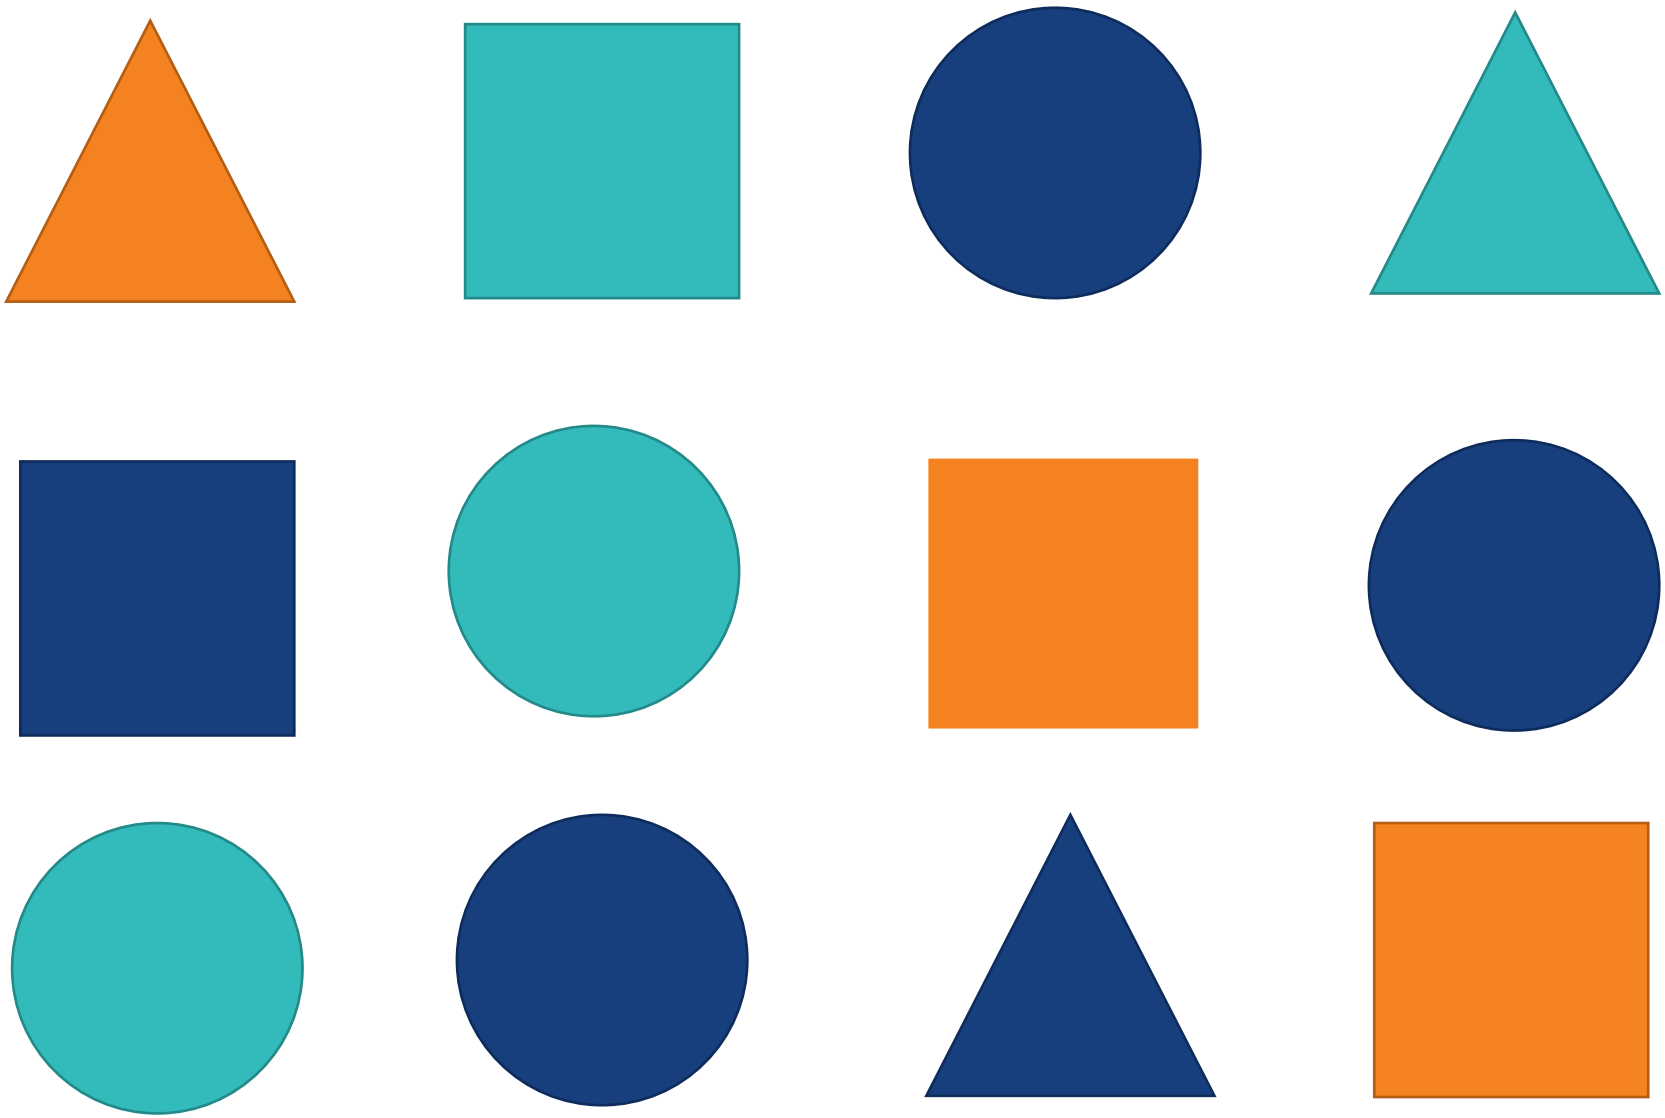

What's the latent variable?

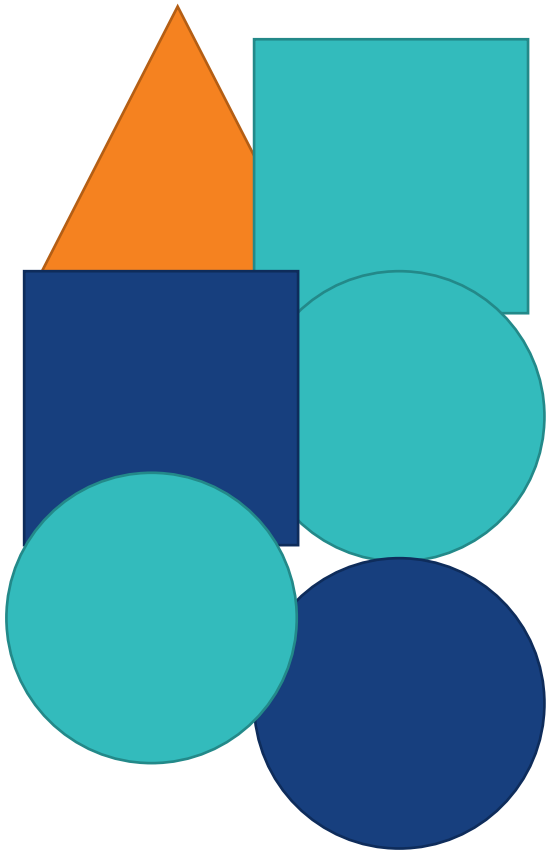

What's the latent variable?

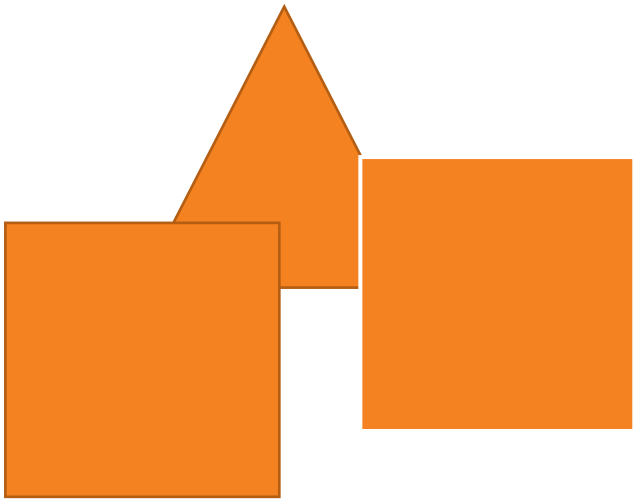

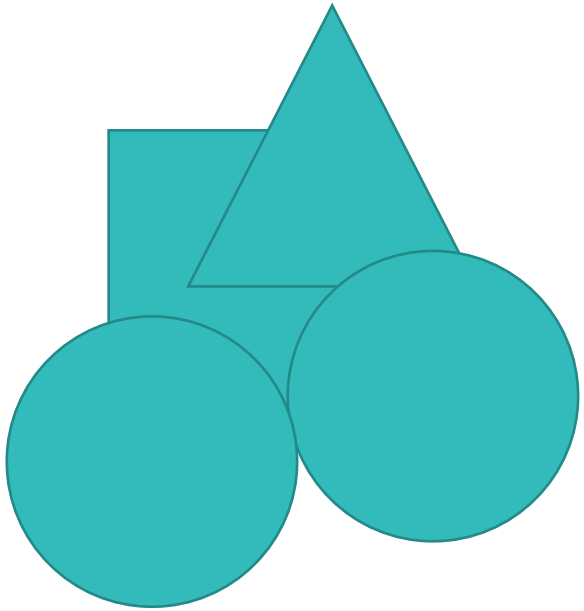

Text Example

- **What's unusual?**
  - Anomaly detection

Date	Amount	Location
2 Mar	$40	Penang
5 Mar	$20	KL
17 Mar	$30	KL
4 Apr	$80	Ipoh
9 Apr	$30	KL
14 Apr	$70	KL
20 May	$100	Johor
25 May	$20	KL
31 May	$3	Kiev
4 Jun	$40	KL
23 Jun	$50	KL
30 Jun	$30	KL
16 Jul	$70	Ipoh
16 Jul	$50	Ipoh

DOCUMENT CLUSTERING

What is text clustering?

- **Clustering is the task of grouping a set of documents in such a way that the documents in each group are more "similar" to each other than to documents in other groups.**
- **Clustering lets you explore your data**
  - Many tools are interactive
- **You can understand your data better, e.g.:**
  - What groupings exist in your data?
  - How many are there? How big is each group?
  - What are the common terms?
  - Are there anomalies?

Clustering Example

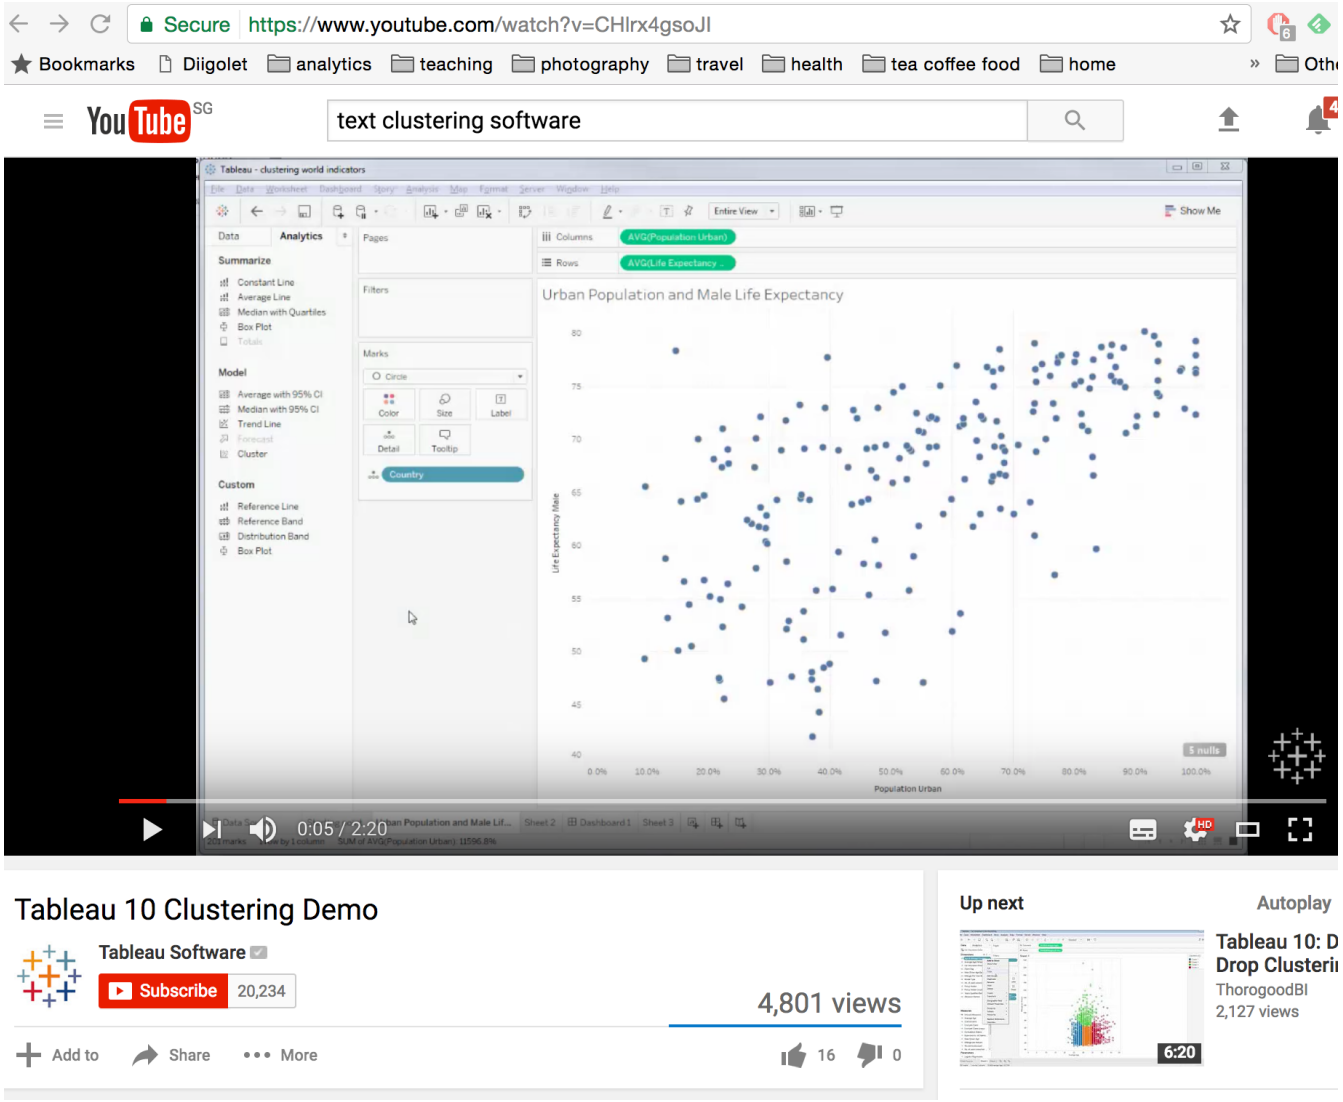

From: https://www.youtube.com/watch?v=CHlrx4gsoJI

Patent Clustering

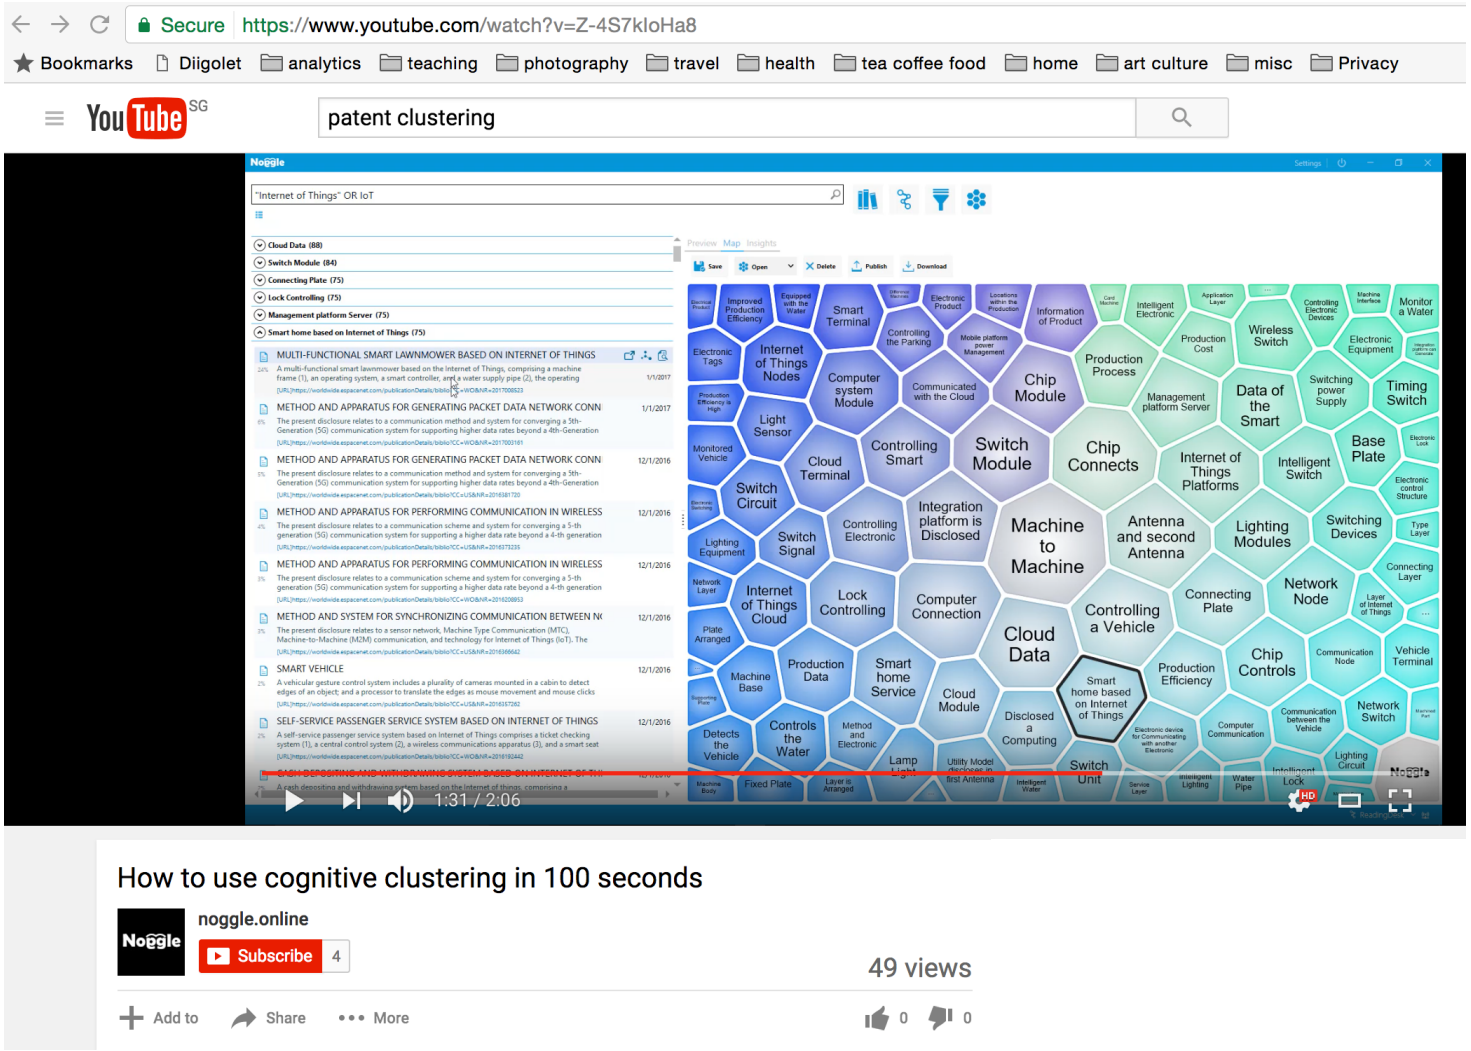

From: https://www.youtube.com/watch?v=Z-4S7kloHa8

Notes about Clustering

- **Your clusters may surprise you**
  - Documents tend to fall into natural classes (clusters)
  - There will be some surprising ones (worth drilling down!)
- **You can control the number of clusters (depends on the algorithm)**
  - You don't want too many clusters (overfit!)
  - You don't want too few clusters (meaningless)
  - **Clusters should lead to fulfilling business outcomes**
- **You don't need training phase to create clusters**
  - Clustering can be language independent (but monolingual)

TOPIC MODELING

What is topic modeling?

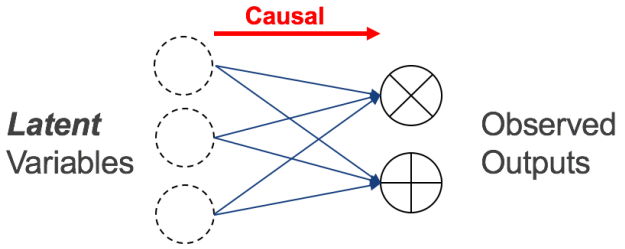

- **Recall…**
  - Latent (or hidden) variables?
  - Output are collections of documents
- **"Topic" modeling**
  - Can we figure out what discourses (==latent variables) would generate the collection of documents?
  - These discourses are just bunches of words
    - If done well, the bunches of words would seem naturally to be together, e.g.,
      - "wag", "bark", "bone", "bite", "dog"
      - "pilot", "plane", "wing", "flight"
  - These bunches of words constitute **topics**

Animation of topic modeling

- Columns = documents
- Rows = words
- Squares = frequency
- Darker = higher frequency

- Group:
  - Documents using similar words
  - Words which occur in similar documents

Resulting set of words are "topics"

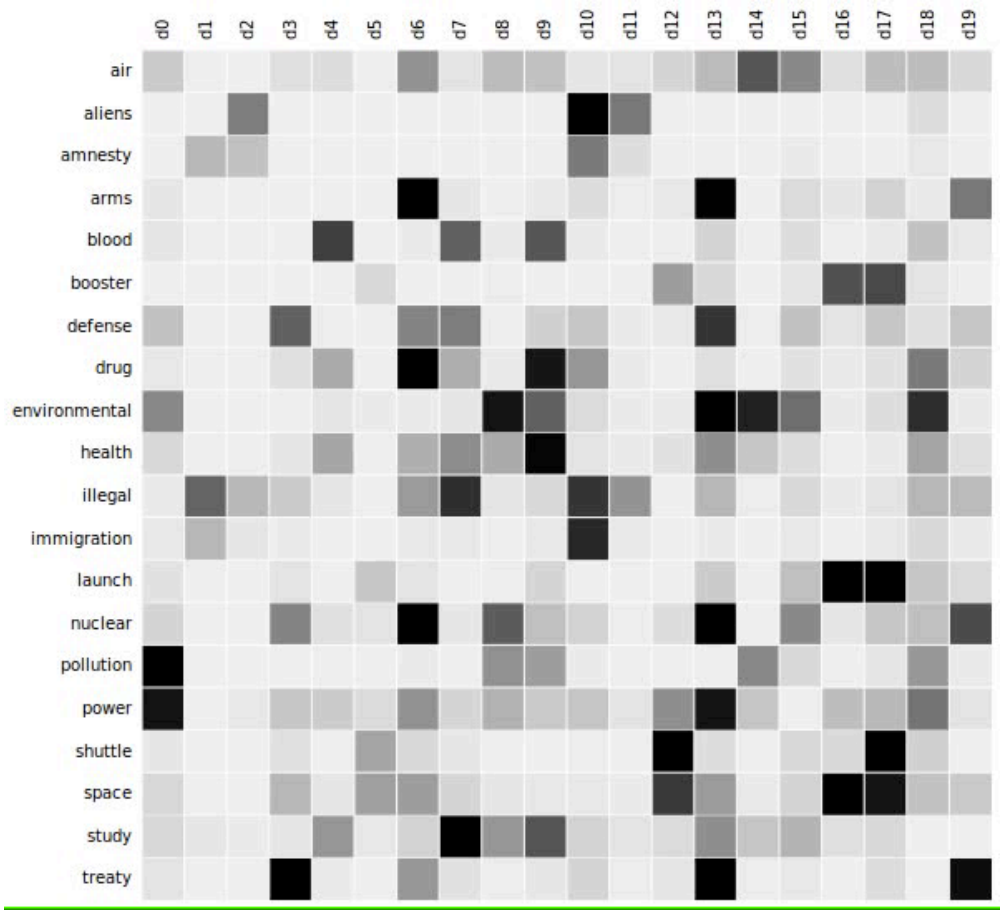

From: http://topicmodels.west.uni-koblenz.de/ckling/tmt/svd_ap.html

LDA Topic Model Explanation

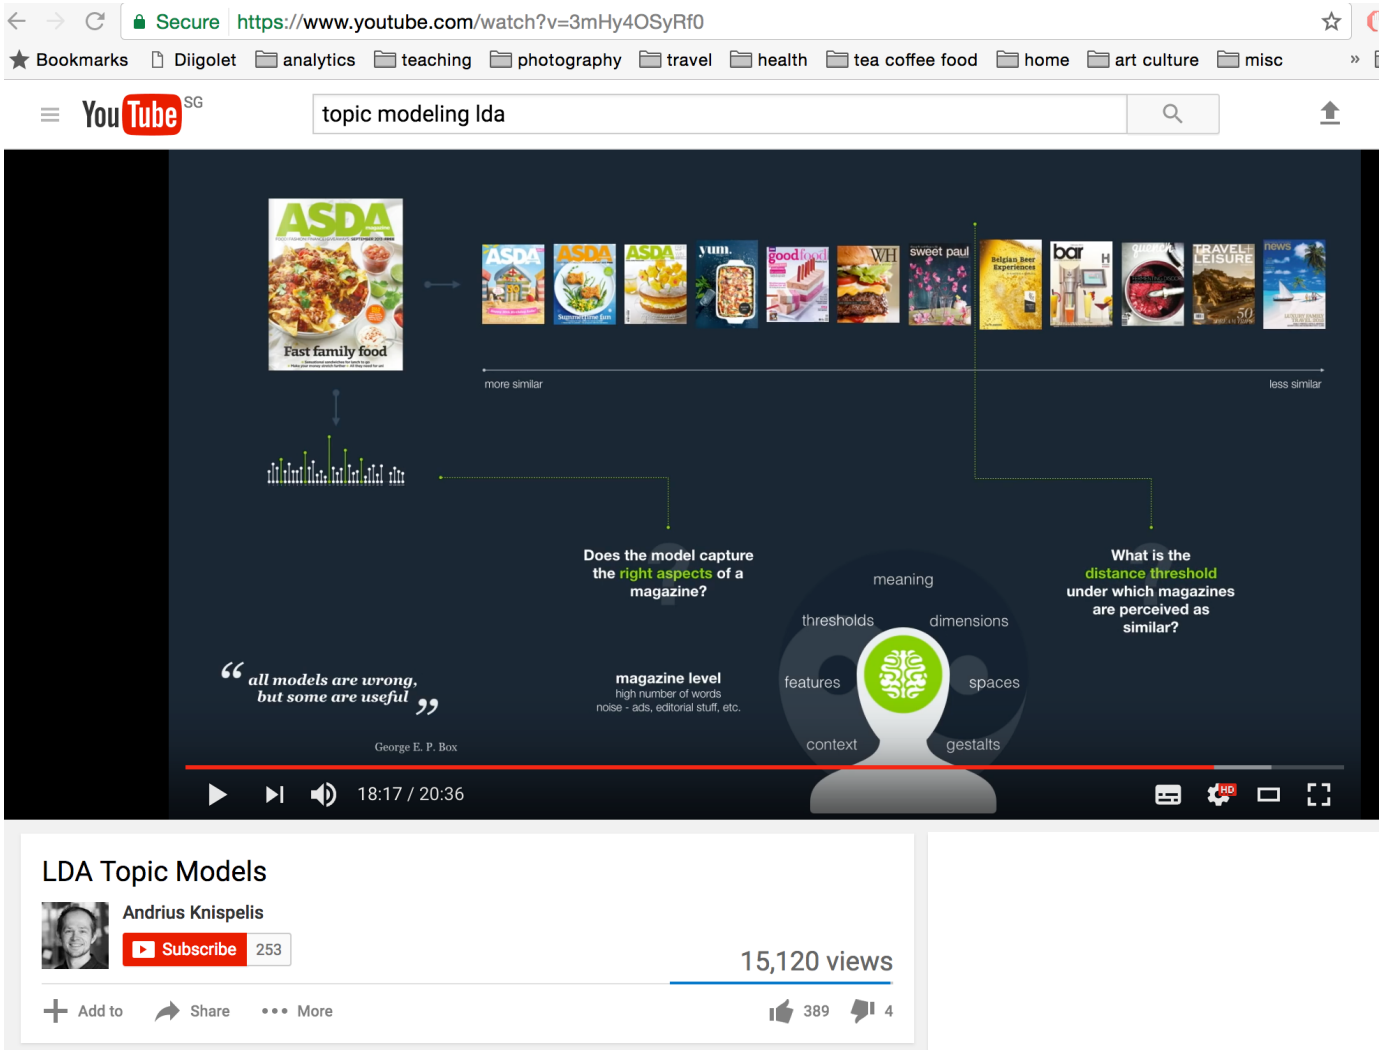

From: https://www.youtube.com/watch?v=3mHy4OSyRf0

More examples of applications

- **Analysis of text, e.g.,**
  - Diachronic analysis:
    - Speeches during election campaign
      - Economy, abortion, build wall, reduce taxes,…
    - Speeches after taking office
      - Reduce taxes, create jobs, immigration, China,…
  - Contrast analysis:
    - Different candidates positions and issues
    - Characteristics of various media publications
  - Evolution of discourse over time and space
    - Word usage and diffusion
    - Engagement patterns

Reference & Resources

- Fabrizio Sebastiani, *A Tutorial on Automated Text Categorization*, web.iiit.ac.in/~jawahar/PRA-03/textCat.pdf
- F. Aiolli, *Text Categorization*, downloaded from http://www.math.unipd.it/~aiolli/corsi/SI-0607/Lez09.251006.pdf
- John Elder, Gary Miner, Bob Nisbet. *Practical Text Mining and Statistical Analysis for non-Structured Text Data Applications*, Academic Press, 2012
- Chris Manning & Hinrich Schutze, *Foundations of Statistical Natural Language Processing*, MIT Press, 1999
- Scott Weingart, *Topic Modeling for Humanists: A Guided Tour*, downloaded from http://www.scottbot.net/HIAL/index.html@p=19113.html
- Ted Underwood, *Topic Modeling made just simple enough*, downloaded from https://tedunderwood.com/2012/04/07/topic-modeling-made-just-simple-enough/
- NLP resources: http://nlp.stanford.edu/links/statnlp.html